

EDA_Assg_NYC_Taxi_Starter

June 25, 2025

1 New York City Yellow Taxi Data

1.1 Objective

In this case study you will be learning exploratory data analysis (EDA) with the help of a dataset on yellow taxi rides in New York City. This will enable you to understand why EDA is an important step in the process of data science and machine learning.

1.2 Problem Statement

As an analyst at an upcoming taxi operation in NYC, you are tasked to use the 2023 taxi trip data to uncover insights that could help optimise taxi operations. The goal is to analyse patterns in the data that can inform strategic decisions to improve service efficiency, maximise revenue, and enhance passenger experience.

1.3 Tasks

You need to perform the following steps for successfully completing this assignment: 1. Data Loading 2. Data Cleaning 3. Exploratory Analysis: Bivariate and Multivariate 4. Creating Visualisations to Support the Analysis 5. Deriving Insights and Stating Conclusions

NOTE: The marks given along with headings and sub-headings are cumulative marks for those particular headings/sub-headings.

The actual marks for each task are specified within the tasks themselves.

For example, marks given with heading 2 or sub-heading 2.1 are the cumulative marks, for your reference only.

The marks you will receive for completing tasks are given with the tasks.

Suppose the marks for two tasks are: 3 marks for 2.1.1 and 2 marks for 3.2.2, or * 2.1.1 [3 marks]
* 3.2.2 [2 marks]

then, you will earn 3 marks for completing task 2.1.1 and 2 marks for completing task 3.2.2.

1.4 Data Understanding

The yellow taxi trip records include fields capturing pick-up and drop-off dates/times, pick-up and drop-off locations, trip distances, itemized fares, rate types, payment types, and driver-reported passenger counts.

The data is stored in Parquet format (*.parquet*). The dataset is from 2009 to 2024. However, for this assignment, we will only be using the data from 2023.

The data for each month is present in a different parquet file. You will get twelve files for each of the months in 2023.

The data was collected and provided to the NYC Taxi and Limousine Commission (TLC) by technology providers like vendors and taxi hailing apps.

You can find the link to the TLC trip records page here: <https://www.nyc.gov/site/tlc/about/tlc-trip-record-data.page>

1.4.1 Data Description

You can find the data description here: [Data Dictionary](#)

Trip Records

Field Name	description
VendorID	A code indicating the TPEP provider that provided the record. 1= Creative Mobile Technologies, LLC; 2= VeriFone Inc.
tpep_pickup_datetime	The date and time when the meter was engaged.
tpep_dropoff_datetime	The date and time when the meter was disengaged.
Passenger_count	The number of passengers in the vehicle. This is a driver-entered value.
Trip_distance	The elapsed trip distance in miles reported by the taximeter.
PULocationID	TLC Taxi Zone in which the taximeter was engaged
DOLocationID	TLC Taxi Zone in which the taximeter was disengaged
RateCodeID	The final rate code in effect at the end of the trip. 1 = Standard rate 2 = JFK 3 = Newark 4 = Nassau or Westchester 5 = Negotiated fare 6 = Group ride

Field Name	description
Store_and_fwd_flag	This flag indicates whether the trip record was held in vehicle memory before sending to the vendor, aka “store and forward,” because the vehicle did not have a connection to the server. Y= store and forward trip N= not a store and forward trip
Payment_type	A numeric code signifying how the passenger paid for the trip. 1 = Credit card 2 = Cash 3 = No charge 4 = Dispute 5 = Unknown 6 = Voided trip
Fare_amount	The time-and-distance fare calculated by the meter. Extra Miscellaneous extras and surcharges. Currently, this only includes the 0.50 and 1 USD rush hour and overnight charges.
MTA_tax	0.50 USD MTA tax that is automatically triggered based on the metered rate in use.
Improvement_surcharge	0.30 USD improvement surcharge assessed trips at the flag drop. The improvement surcharge began being levied in 2015.
Tip_amount	Tip amount – This field is automatically populated for credit card tips. Cash tips are not included.
Tolls_amount	Total amount of all tolls paid in trip.
total_amount	The total amount charged to passengers. Does not include cash tips.
Congestion_Surcharge	Total amount collected in trip for NYS congestion surcharge.
Airport_fee	1.25 USD for pick up only at LaGuardia and John F. Kennedy Airports

Although the amounts of extra charges and taxes applied are specified in the data dictionary, you will see that some cases have different values of these charges in the actual data.

Taxi Zones

Each of the trip records contains a field corresponding to the location of the pickup or drop-off of the trip, populated by numbers ranging from 1-263.

These numbers correspond to taxi zones, which may be downloaded as a table or map/shapefile and matched to the trip records using a join.

This is covered in more detail in later sections.

1.5 1 Data Preparation

[5 marks]

1.5.1 Import Libraries

```
[140]: # Import warnings
import warnings
```

```
[142]: # Import the libraries you will be using for analysis
import numpy as np
import pandas as pd
import matplotlib as mlib
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[144]: # Recommended versions
# numpy version: 1.26.4
# pandas version: 2.2.2
# matplotlib version: 3.10.0
# seaborn version: 0.13.2

# Check versions
print("numpy version:", np.__version__)
print("pandas version:", pd.__version__)
print("matplotlib version:", plt.matplotlib.__version__)
print("seaborn version:", sns.__version__)
```

```
numpy version: 1.26.4
pandas version: 2.2.3
matplotlib version: 3.10.0
seaborn version: 0.13.2
```

1.5.2 1.1 Load the dataset

[5 marks]

You will see twelve files, one for each month.

To read parquet files with Pandas, you have to follow a similar syntax as that for CSV files.

```
df = pd.read_parquet('file.parquet')
```

```
[148]: # Try loading one file
```

```
df = pd.read_parquet('2023-1.parquet')
# df = pd.read_parquet('2023-1.parquet')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3041714 entries, 0 to 3066765
Data columns (total 19 columns):
#   Column                Dtype
---  -
0   VendorID              int64
1   tpep_pickup_datetime  datetime64[us]
2   tpep_dropoff_datetime datetime64[us]
3   passenger_count       float64
4   trip_distance         float64
5   RatecodeID            float64
6   store_and_fwd_flag    object
7   PULocationID          int64
8   DOLocationID          int64
9   payment_type          int64
10  fare_amount           float64
11  extra                 float64
12  mta_tax               float64
13  tip_amount            float64
14  tolls_amount          float64
15  improvement_surcharge float64
16  total_amount          float64
17  congestion_surcharge  float64
18  airport_fee           float64
dtypes: datetime64[us](2), float64(12), int64(4), object(1)
memory usage: 464.1+ MB
```

```
[150]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3041714 entries, 0 to 3066765
Data columns (total 19 columns):
#   Column                Dtype
---  -
0   VendorID              int64
1   tpep_pickup_datetime  datetime64[us]
2   tpep_dropoff_datetime datetime64[us]
3   passenger_count       float64
4   trip_distance         float64
5   RatecodeID            float64
6   store_and_fwd_flag    object
7   PULocationID          int64
8   DOLocationID          int64
9   payment_type          int64
```

```

10 fare_amount          float64
11 extra                float64
12 mta_tax              float64
13 tip_amount           float64
14 tolls_amount         float64
15 improvement_surcharge float64
16 total_amount         float64
17 congestion_surcharge float64
18 airport_fee          float64
dtypes: datetime64[us](2), float64(12), int64(4), object(1)
memory usage: 464.1+ MB

```

How many rows are there? Do you think handling such a large number of rows is computationally feasible when we have to combine the data for all twelve months into one?

To handle this, we need to sample a fraction of data from each of the files. How to go about that? Think of a way to select only some portion of the data from each month's file that accurately represents the trends.

Sampling the Data

One way is to take a small percentage of entries for pickup in every hour of a date. So, for all the days in a month, we can iterate through the hours and select 5% values randomly from those. Use `tpep_pickup_datetime` for this. Separate date and hour from the datetime values and then for each date, select some fraction of trips for each of the 24 hours.

To sample data, you can use the `sample()` method. Follow this syntax:

```

# sampled_data is an empty DF to keep appending sampled data of each hour
# hour_data is the DF of entries for an hour 'X' on a date 'Y'

sample = hour_data.sample(frac = 0.05, random_state = 42)
# sample 0.05 of the hour_data
# random_state is just a seed for sampling, you can define it yourself

sampled_data = pd.concat([sampled_data, sample]) # adding data for this hour to the DF

```

This `sampled_data` will contain 5% values selected at random from each hour.

Note that the code given above is only the part that will be used for sampling and not the complete code required for sampling and combining the data files.

Keep in mind that you sample by date AND hour, not just hour. (Why?)

1.1.1 [5 marks] Figure out how to sample and combine the files.

Note: It is not mandatory to use the method specified above. While sampling, you only need to make sure that your sampled data represents the overall data of all the months accurately.

```
[152]: # Sample the data
# It is recommended to not load all the files at once to avoid memory overload
```

```
def extractMonthlySampleData (filePath):
    df = pd.read_parquet(filePath)  #'trip_records/2023-1.parquet'
    # Extract date and hour from pickup datetime
    df['pickup_date'] = df['tpep_pickup_datetime'].dt.date
    df['pickup_hour'] = df['tpep_pickup_datetime'].dt.hour
    df_group_day_hour = df.groupby(["pickup_date", "pickup_hour"])
    sample_df_rows = []

    for (pickup_date, pickup_hour), hour_data in df_group_day_hour:
        sampled = hour_data.sample(frac=0.05, random_state=42)
        sample_df_rows.append(sampled)

    monthly_sample = pd.concat(sample_df_rows, ignore_index=True)
    return monthly_sample
```

```
[156]: sample = extractMonthlySampleData ('2023-1.parquet')
sample.head()
```

```
[156]: VendorID tpep_pickup_datetime tpep_dropoff_datetime passenger_count \
0          2  2022-12-31 23:51:30    2022-12-31 23:56:06              1.0
1          2  2023-01-01 00:07:18    2023-01-01 00:23:15              1.0
2          2  2023-01-01 00:16:41    2023-01-01 00:21:46              2.0
3          2  2023-01-01 00:14:03    2023-01-01 00:24:36              3.0
4          2  2023-01-01 00:24:30    2023-01-01 00:29:55              1.0

      trip_distance  RatecodeID store_and_fwd_flag  PULocationID  DOLocationID \
0              0.86          1.0                N           141           140
1              7.74          1.0                N           138           256
2              1.24          1.0                N           161           237
3              1.44          1.0                N           237           141
4              0.54          1.0                N           143           142

      payment_type  ...  extra  mta_tax  tip_amount  tolls_amount \
0              1  ...    1.0     0.5         2.00         0.0
1              2  ...    6.0     0.5         0.00         0.0
2              1  ...    1.0     0.5         2.58         0.0
3              2  ...    1.0     0.5         0.00         0.0
4              2  ...    1.0     0.5         0.00         0.0

      improvement_surcharge  total_amount  congestion_surcharge  airport_fee \
0              1.0              13.50              2.5          0.00
1              1.0              41.15              0.0          1.25
```

2	1.0	15.48	2.5	0.00
3	1.0	16.40	2.5	0.00
4	1.0	11.50	2.5	0.00

	pickup_date	pickup_hour
0	2022-12-31	23
1	2023-01-01	0
2	2023-01-01	0
3	2023-01-01	0
4	2023-01-01	0

[5 rows x 21 columns]

```
[158]: # from google.colab import drive
# drive.mount('/content/drive')
print(f"sample : {sample.shape}")
```

sample : (152087, 21)

```
[160]: # Take a small percentage of entries from each hour of every date.
# Iterating through the monthly data:
#   read a month file -> day -> hour: append sampled data -> move to next hour
#   ↳-> move to next day after 24 hours -> move to next month file
# Create a single dataframe for the year combining all the monthly data

# Select the folder having data files
import os

# Select the folder having data files
os.chdir(r'C:\Users\Sherry\UpgradExample\Datasets and Dictionary-NYC\Starter_
↳Notebook - EDA NYC Taxi\trip_records')

# Create a list of all the twelve files to read
file_list = os.listdir()

# initialise an empty dataframe
df = pd.DataFrame()

# iterate through the list of files and sample one by one:
for file_name in file_list:
    try:
        # file path for the current file
        file_path = os.path.join(os.getcwd(), file_name)

        # Reading the current file
```



```

        # We will store the sampled data for the current date in this df by
        ↳ appending the sampled data from each hour to this
        # After completing iteration through each date, we will append this
        ↳ data to the final dataframe.
        sampled_data = pd.DataFrame()
        sampled_data = extractMonthlySampleData(file_path)
        print (f"file_name : {file_name}  sampled_data.shape:{sampled_data.
        ↳ shape}")
        # Loop through dates and then loop through every hour of each date

        # Iterate through each hour of the selected date

        # Sample 5% of the hourly data randomly

        # add data of this hour to the dataframe

        # Concatenate the sampled data of all the dates to a single dataframe
        df = pd.concat([df,sampled_data])
        print (f"file_name : {file_name}  df.shape:{df.shape}")
        #df = # we initialised this empty DF earlier

    except Exception as e:
        print(f"Error reading file {file_name}: {e}")

print("completed processing")
#df = pd.concat(df, ignore_index=True)
df.info()

```

```

file_name : 2023-1.parquet  sampled_data.shape:(152087, 21)
file_name : 2023-1.parquet  df.shape:(152087, 21)
file_name : 2023-10.parquet  sampled_data.shape:(174255, 21)
file_name : 2023-10.parquet  df.shape:(326342, 22)
file_name : 2023-11.parquet  sampled_data.shape:(165133, 21)
file_name : 2023-11.parquet  df.shape:(491475, 22)
file_name : 2023-12.parquet  sampled_data.shape:(166709, 21)
file_name : 2023-12.parquet  df.shape:(658184, 22)
file_name : 2023-2.parquet  sampled_data.shape:(168696, 21)
file_name : 2023-2.parquet  df.shape:(826880, 22)
file_name : 2023-3.parquet  sampled_data.shape:(163786, 21)
file_name : 2023-3.parquet  df.shape:(990666, 22)
file_name : 2023-4.parquet  sampled_data.shape:(139641, 21)
file_name : 2023-4.parquet  df.shape:(1130307, 22)
file_name : 2023-5.parquet  sampled_data.shape:(144458, 21)
file_name : 2023-5.parquet  df.shape:(1274765, 22)
file_name : 2023-6.parquet  sampled_data.shape:(162910, 21)
file_name : 2023-6.parquet  df.shape:(1437675, 22)

```

```

file_name : 2023-7.parquet  sampled_data.shape:(174068, 21)
file_name : 2023-7.parquet  df.shape:(1611743, 22)
file_name : 2023-8.parquet  sampled_data.shape:(143782, 21)
file_name : 2023-8.parquet  df.shape:(1755525, 22)
file_name : 2023-9.parquet  sampled_data.shape:(140875, 21)
file_name : 2023-9.parquet  df.shape:(1896400, 22)
file_name : sample-2023.parquet  sampled_data.shape:(94927, 22)
file_name : sample-2023.parquet  df.shape:(1991327, 22)
completed processing
<class 'pandas.core.frame.DataFrame'>
Index: 1991327 entries, 0 to 94926
Data columns (total 22 columns):
 #   Column                Dtype
---  -
 0   VendorID              int64
 1   tpep_pickup_datetime  datetime64[us]
 2   tpep_dropoff_datetime datetime64[us]
 3   passenger_count       float64
 4   trip_distance         float64
 5   RatecodeID            float64
 6   store_and_fwd_flag    object
 7   PULocationID          int64
 8   DOLocationID          int64
 9   payment_type          int64
10   fare_amount           float64
11   extra                 float64
12   mta_tax               float64
13   tip_amount            float64
14   tolls_amount          float64
15   improvement_surcharge float64
16   total_amount          float64
17   congestion_surcharge  float64
18   airport_fee           float64
19   pickup_date           object
20   pickup_hour           int32
21   Airport_fee           float64
dtypes: datetime64[us](2), float64(13), int32(1), int64(4), object(2)
memory usage: 341.8+ MB

```

After combining the data files into one DataFrame, convert the new DataFrame to a CSV or parquet file and store it to use directly.

Ideally, you can try keeping the total entries to around 250,000 to 300,000.

```
[178]: df.to_csv('sample-2023-ex1.csv', index=False)
```

```
[162]: # Store the df in csv/parquet
df.to_parquet('sample-2023-v1.parquet')
```

```
print (df.shape)
```

```
(1991327, 22)
```

1.6 2 Data Cleaning

[30 marks]

Now we can load the new data directly.

```
[180]: # Load the new data file
```

```
df = pd.read_parquet('sample-2023-v1.parquet')
```

```
[182]: df.head()
```

```
[182]:
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	\
0	2	2022-12-31 23:51:30	2022-12-31 23:56:06	1.0	
1	2	2023-01-01 00:07:18	2023-01-01 00:23:15	1.0	
2	2	2023-01-01 00:16:41	2023-01-01 00:21:46	2.0	
3	2	2023-01-01 00:14:03	2023-01-01 00:24:36	3.0	
4	2	2023-01-01 00:24:30	2023-01-01 00:29:55	1.0	

	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	\
0	0.86	1.0	N	141	140	
1	7.74	1.0	N	138	256	
2	1.24	1.0	N	161	237	
3	1.44	1.0	N	237	141	
4	0.54	1.0	N	143	142	

	payment_type	...	mta_tax	tip_amount	tolls_amount	\
0	1	...	0.5	2.00	0.0	
1	2	...	0.5	0.00	0.0	
2	1	...	0.5	2.58	0.0	
3	2	...	0.5	0.00	0.0	
4	2	...	0.5	0.00	0.0	

	improvement_surcharge	total_amount	congestion_surcharge	airport_fee	\
0	1.0	13.50	2.5	0.00	
1	1.0	41.15	0.0	1.25	
2	1.0	15.48	2.5	0.00	
3	1.0	16.40	2.5	0.00	
4	1.0	11.50	2.5	0.00	

	pickup_date	pickup_hour	Airport_fee
0	2022-12-31	23	NaN
1	2023-01-01	0	NaN
2	2023-01-01	0	NaN

```

3    2023-01-01          0      NaN
4    2023-01-01          0      NaN

```

[5 rows x 22 columns]

[184]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
Index: 1991327 entries, 0 to 94926
Data columns (total 22 columns):
 #   Column              Dtype
---  -
 0   VendorID            int64
 1   tpep_pickup_datetime  datetime64[us]
 2   tpep_dropoff_datetime datetime64[us]
 3   passenger_count      float64
 4   trip_distance        float64
 5   RatecodeID          float64
 6   store_and_fwd_flag   object
 7   PULocationID        int64
 8   DOLocationID        int64
 9   payment_type        int64
10   fare_amount         float64
11   extra               float64
12   mta_tax             float64
13   tip_amount          float64
14   tolls_amount        float64
15   improvement_surcharge float64
16   total_amount        float64
17   congestion_surcharge float64
18   airport_fee         float64
19   pickup_date         object
20   pickup_hour         int32
21   Airport_fee         float64
dtypes: datetime64[us](2), float64(13), int32(1), int64(4), object(2)
memory usage: 341.8+ MB

```

2.1 Fixing Columns [10 marks]

Fix/drop any columns as you seem necessary in the below sections

2.1.1 [2 marks]

Fix the index and drop unnecessary columns

[186]: *# Fix the index and drop any columns that are not needed*

```

df.reset_index(drop=True, inplace=True)
# All columns provide some information which can be useful

```

2.1.2 [3 marks] There are two airport fee columns. This is possibly an error in naming columns. Let's see whether these can be combined into a single column.

```
[188]: # Combine the two airport fee columns
print (df[(df['Airport_fee'].isnull()) & (df['airport_fee'].isnull())].shape)
print (df[(df['Airport_fee'].isnull()) & (df['airport_fee'].notnull())].shape)
print (df[(df['Airport_fee'].notnull()) & (df['airport_fee'].isnull())].shape)
print (df[(df['Airport_fee'].notnull()) & (df['airport_fee'].notnull())].shape)

df['airport_fee'] = df['airport_fee'].fillna(df['Airport_fee'])
print (df[(df['Airport_fee'].isnull()) & (df['airport_fee'].notnull())].shape)
# (Optional) Drop the duplicate column after merging
df.drop(columns='Airport_fee', inplace=True)

print(df['airport_fee'].isnull().sum())
```

```
(68132, 22)
(155904, 22)
(1767291, 22)
(0, 22)
(155904, 22)
68132
```

2.1.3 [5 marks] Fix columns with negative (monetary) values

```
[190]: # check where values of fare amount are negative

negative_fares = df[df['fare_amount'] < 0]
negative_fares.head()
print (negative_fares.shape)
```

```
(0, 21)
```

Did you notice something different in the RatecodeID column for above records?

```
[ ]: # Analyse RatecodeID for the negative fare amounts
## I have created sample multiple times but no fare amount was negative, though
↳ there is total amount as negative but that is because of other negative
↳ components
```

```
[210]: # Find which columns have negative values

numeric_df = df.select_dtypes(include='number')
negative_columns = numeric_df.columns[(numeric_df < 0).any()]

print("Columns with negative values:", list(negative_columns))
```

```
Columns with negative values: ['extra', 'mta_tax', 'improvement_surcharge',
'total_amount', 'congestion_surcharge', 'airport_fee']
```

```
[224]: # fix these negative values
# will use abs function to make negative to positive numbers
(df[negative_columns] < 0).sum()
df[negative_columns] = df[negative_columns].abs()
(df[negative_columns] < 0).sum()
```

```
[224]: extra          0
mta_tax            0
improvement_surcharge  0
total_amount       0
congestion_surcharge  0
airport_fee        0
dtype: int64
```

1.6.1 2.2 Handling Missing Values

[10 marks]

2.2.1 [2 marks] Find the proportion of missing values in each column

```
[240]: # Find the proportion of missing values in each column

print(f"proportion of values on percentage of records: \n {df.isnull().mean() * 100} ")
print(f"proportion of values on number of records: \n {df.isnull().sum()} ")
```

proportion of values on percentage of records:

```
VendorID          0.000000
tpep_pickup_datetime  0.000000
tpep_dropoff_datetime  0.000000
passenger_count     3.421437
trip_distance       0.000000
RatecodeID         3.421437
store_and_fwd_flag   3.421437
PULocationID       0.000000
DOLocationID       0.000000
payment_type       0.000000
fare_amount        0.000000
extra              0.000000
mta_tax            0.000000
tip_amount         0.000000
tolls_amount       0.000000
improvement_surcharge  0.000000
total_amount       0.000000
congestion_surcharge  3.421437
airport_fee        3.421437
pickup_date        0.000000
pickup_hour        0.000000
dtype: float64
```

proportion of values on number of records:

```
VendorID          0
tpep_pickup_datetime  0
tpep_dropoff_datetime  0
passenger_count    68132
trip_distance      0
RatecodeID        68132
store_and_fwd_flag  68132
PULocationID      0
DOLocationID      0
payment_type       0
fare_amount        0
extra              0
mta_tax            0
tip_amount         0
tolls_amount       0
improvement_surcharge  0
total_amount       0
congestion_surcharge  68132
airport_fee        68132
pickup_date        0
pickup_hour        0
dtype: int64
```

2.2.2 [3 marks] Handling missing values in passenger_count

```
[266]: # Display the rows with null values
null_rows = df[df.isnull().any(axis=1)]
print("null_rows Shape: ", null_rows.shape)
print(null_rows.head(10))

# Impute NaN values in 'passenger_count'
most_commonly_used_value = df['passenger_count'].mode()[0]
print("most_commonly_used_value :",most_commonly_used_value)

print ("Before - number of nulls in passenger_count",df['passenger_count'].
      ↪isnull().sum())
df['passenger_count'].fillna(most_commonly_used_value, inplace=True)

print ("After setting the most_commonly_used_value - number of nulls in_
      ↪passenger_count",df['passenger_count'].isnull().sum())
```

null_rows Shape: (68132, 21)

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	\
5	2	2023-01-01 00:43:00	2023-01-01 01:01:00	NaN	
16	2	2023-01-01 00:41:50	2023-01-01 01:14:50	NaN	
43	2	2023-01-01 00:37:21	2023-01-01 00:54:18	NaN	

44	2	2023-01-01 00:44:03	2023-01-01 01:13:49	NaN
47	2	2023-01-01 00:50:55	2023-01-01 01:19:06	NaN
53	2	2023-01-01 00:55:00	2023-01-01 01:07:00	NaN
70	1	2023-01-01 00:28:22	2023-01-01 00:41:25	NaN
78	1	2023-01-01 00:37:09	2023-01-01 00:58:16	NaN
105	2	2023-01-01 00:58:50	2023-01-01 01:17:07	NaN
126	2	2023-01-01 00:16:00	2023-01-01 00:38:00	NaN

	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	\
5	19.24	NaN	None	66	107	
16	10.77	NaN	None	151	106	
43	4.52	NaN	None	114	262	
44	9.19	NaN	None	239	256	
47	2.74	NaN	None	90	48	
53	1.89	NaN	None	141	143	
70	2.90	NaN	None	263	137	
78	0.00	NaN	None	36	7	
105	4.43	NaN	None	33	68	
126	4.77	NaN	None	90	262	

	payment_type	...	extra	mta_tax	tip_amount	tolls_amount	\
5	0	...	0.0	0.5	5.93	0.00	
16	0	...	0.0	0.5	11.19	6.55	
43	0	...	0.0	0.5	0.00	0.00	
44	0	...	0.0	0.5	2.20	0.00	
47	0	...	0.0	0.5	3.37	0.00	
53	0	...	0.0	0.5	3.36	0.00	
70	0	...	1.0	0.5	3.09	0.00	
78	0	...	1.0	0.5	0.00	0.00	
105	0	...	0.0	0.5	4.37	0.00	
126	0	...	0.0	0.5	7.86	0.00	

	improvement_surcharge	total_amount	congestion_surcharge	airport_fee	\
5	1.0	35.57	NaN	NaN	
16	1.0	67.12	NaN	NaN	
43	1.0	29.38	NaN	NaN	
44	1.0	46.20	NaN	NaN	
47	1.0	25.85	NaN	NaN	
53	1.0	20.16	NaN	NaN	
70	1.0	23.69	NaN	NaN	
78	1.0	32.00	NaN	NaN	
105	1.0	33.50	NaN	NaN	
126	1.0	39.28	NaN	NaN	

	pickup_date	pickup_hour
5	2023-01-01	0
16	2023-01-01	0
43	2023-01-01	0


```

44    2023-01-01          0
47    2023-01-01          0
53    2023-01-01          0
70    2023-01-01          0
78    2023-01-01          0
105   2023-01-01          0
126   2023-01-01          0

```

[10 rows x 21 columns]

most_commonly_used_value : 1.0

Before - number of nulls in passenger_count 68132

After setting the most_commonly_used_value - number of nulls in passenger_count
0

C:\Users\Sherry\AppData\Local\Temp\ipykernel_11388\419987922.py:12:

FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['passenger_count'].fillna(most_commonly_used_value, inplace=True)
```

Did you find zeroes in passenger_count? Handle these.

2.2.3 [2 marks] Handle missing values in RatecodeID

```

[278]: # Fix missing values in 'RatecodeID'

most_commonly_used_RatecodeID = 99
print("most_commonly_used_RatecodeID :",most_commonly_used_RatecodeID)

print ("Before the change- number of nulls in RatecodeID",df['RatecodeID'].
      ↪isnull().sum())
print ("will change the value with 99 which is equivalent to null/unknown")

df['RatecodeID'].fillna(most_commonly_used_RatecodeID, inplace=True)

print ("After setting the most_commonly_used_value - number of nulls in_
      ↪passenger_count",df['passenger_count'].isnull().sum())

```

most_commonly_used_RatecodeID : 99

Before the change- number of nulls in RatecodeID 68132

will change the value with 99 which is equivalent to null/unknown

After setting the most_commonly_used_value - number of nulls in passenger_count
0

2.2.4 [3 marks] Impute NaN in congestion_surcharge

```
[288]: # handle null values in congestion_surcharge
# here i will be setting the value as 2.5 after summing up total
↳ fare_amount, extra, mta_tax, tip_amount, tolls_amount, improvement_surcharge and
↳ then subtracting it with Total. A very good finding was that most of the
↳ congestion charge was missing. It should 2.5 as that is the most commonly
↳ used value.

most_commonly_used_cs = df['congestion_surcharge'].mode()[0]
print(most_commonly_used_cs)

df['total_diff'] = df['total_amount'] - (
    df['fare_amount'] +
    df['extra'] +
    df['mta_tax'] +
    df['tip_amount'] +
    df['tolls_amount'] +
    df['improvement_surcharge']
)
```

2.5

```
[292]: check_cs = (df['congestion_surcharge'].isnull()) & (df['total_diff'] == 2.5)
print(f"Before - Rows available for update: {check_cs.sum()}")
df.loc[
    (df['congestion_surcharge'].isnull()) & (df['total_diff'] == 2.5),
    ['congestion_surcharge', 'airport_fee']
] = [2.5, 0.0]

check_cs = (df['congestion_surcharge'].isnull()) & (df['total_diff'] == 2.5)
print(f"After - Rows availablefor update: {check_cs.sum()}")
```

Before - Rows available for update: 49245

After - Rows availablefor update: 0

Are there missing values in other columns? Did you find NaN values in some other set of columns?
Handle those missing values below.

```
[296]: # Handle any remaining missing values
print("number of nulls with congestion_surcharge", df['congestion_surcharge'].
    ↳ isnull().sum())

initial_rows = len(df)
df = df[df['congestion_surcharge'].notnull()]
print(f"Rows after removal: {len(df)} (Removed {initial_rows - len(df)} rows)")
```

```
print("number of nulls with congestion_surcharge",df['congestion_surcharge'].
↳isnull().sum())
print("number of nulls with airport_fee",df['airport_fee'].isnull().sum())
```

number of nulls with congestion_surcharge 18887
 Rows after removal: 1972440 (Removed 18887 rows)
 number of nulls with congestion_surcharge 0
 number of nulls with airport_fee 0

```
[306]: print(f"Before proportion of values on number of records: \n_
↳{df['store_and_fwd_flag'].isnull().sum()} ")
most_commonly_used_store_and_fwd_flag = df['store_and_fwd_flag'].mode()[0]
# Will set the flag as N, as anyways I dont see much use of these column for_
↳now. Can remove this later if not used,
df['store_and_fwd_flag'].fillna('N', inplace=True)
print(f"After proportion of values on number of records: \n_
↳{df['store_and_fwd_flag'].isnull().sum()} ")
```

Before proportion of values on number of records:
 49245

C:\Users\Sherry\AppData\Local\Temp\ipykernel_11388\3598645282.py:4:
 FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
 through chained assignment using an inplace method.
 The behavior will change in pandas 3.0. This inplace method will never work
 because the intermediate object on which we are setting values always behaves as
 a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using
 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)
 instead, to perform the operation inplace on the original object.

```
df['store_and_fwd_flag'].fillna('N', inplace=True)
C:\Users\Sherry\AppData\Local\Temp\ipykernel_11388\3598645282.py:4:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 df['store_and_fwd_flag'].fillna('N', inplace=True)

After proportion of values on number of records:
 0

1.6.2 2.3 Handling Outliers

[10 marks]

Before we start fixing outliers, let's perform outlier analysis.

```
[316]: # Describe the data and check if there are any potential outliers present
# Check for potential out of place values in various columns


df.describe()

fig, axes = plt.subplots(1, 2, figsize=(12, 6))


# Box plot for trip_distance
sns.boxplot(y=df['trip_distance'], ax=axes[0])
axes[0].set_title("Trip Distance")

# Box plot for fare_amount
sns.boxplot(y=df['fare_amount'], ax=axes[1])
axes[1].set_title("Fare Amount")

plt.tight_layout()
plt.show()
```



```
[319]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='trip_distance', y='fare_amount', data=df, alpha=0.5)
plt.title("Scatter Plot of Trip Distance vs Fare Amount")
plt.xlabel("Trip Distance (miles)")
plt.ylabel("Fare Amount ($)")
plt.grid(True)
plt.show()
```



```
[323]: df['passenger_count'].value_counts()
```

```
[323]: passenger_count
1.0    1495716
2.0     291012
3.0     72376
4.0     40387
0.0     31234
5.0     25004
6.0     16690
8.0         11
7.0          5
9.0          5
Name: count, dtype: int64
```

2.3.1 [10 marks] Based on the above analysis, it seems that some of the outliers are present due to errors in registering the trips. Fix the outliers.

Some points you can look for: - Entries where `trip_distance` is nearly 0 and `fare_amount` is more than 300 - Entries where `trip_distance` and `fare_amount` are 0 but the pickup and dropoff zones are different (both distance and fare should not be zero for different zones) - Entries where `trip_distance` is more than 250 miles. - Entries where `payment_type` is 0 (there is no payment_type 0 defined in the data dictionary)

These are just some suggestions. You can handle outliers in any way you wish, using the insights from above outlier analysis.

How will you fix each of these values? Which ones will you drop and which ones will you replace?
First, let us remove 7+ passenger counts as there are very less instances.

```
[325]: # remove passenger_count > 6
df = df[df['passenger_count'] <= 6]
```

```
[327]: df['passenger_count'].value_counts()
```

```
[327]: passenger_count
1.0    1495716
2.0     291012
3.0      72376
4.0     40387
0.0     31234
5.0     25004
6.0     16690
Name: count, dtype: int64
```

```
[339]: # Continue with outlier handling
# remove engtries with trip_distance is 0 and fare_amount is more than 300
df = df[~((df['trip_distance'] == 0) & (df['fare_amount'] > 300))]
```

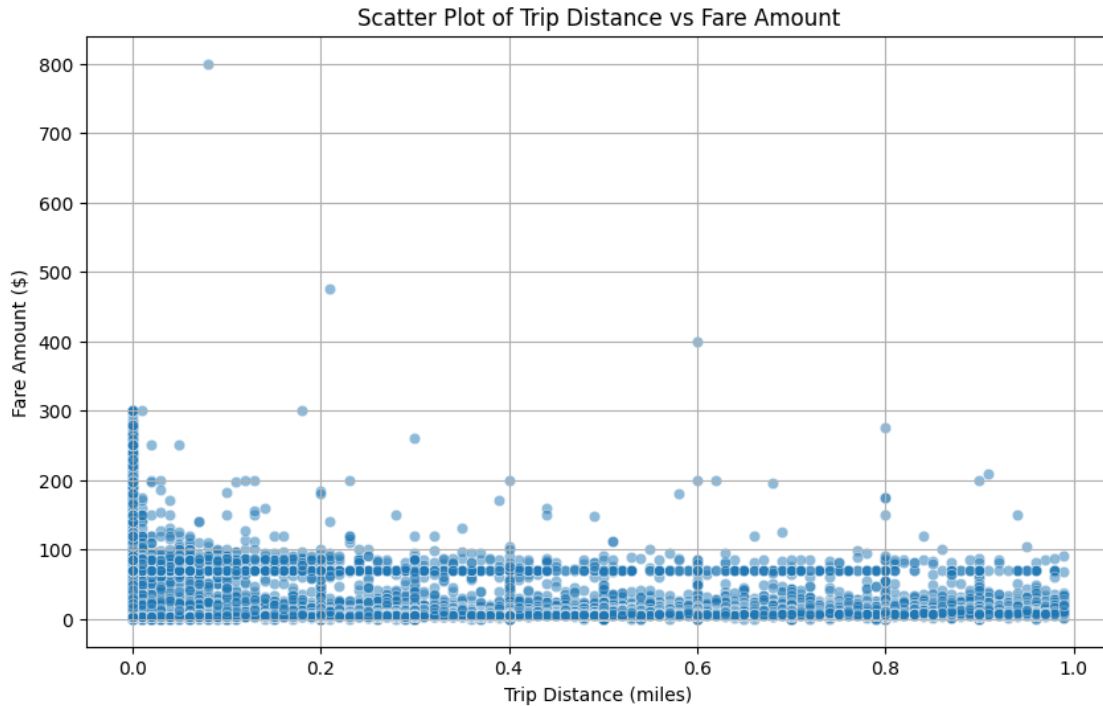
```
[341]: # remove Entries where trip_distance and fare_amount are 0 but the pickup and
↳ dropoff zones are different (both distance and fare should not be zero for
↳ different zones)

df = df[~((df['trip_distance'] == 0) &
          (df['fare_amount'] == 0) &
          (df['PULocationID'] != df['DOLocationID']))]
```

```
[355]: #Remove outlier Fare amounts
print(len( df[~(df['fare_amount'] > 1500)]))
df = df[~(df['fare_amount'] > 1500)]
print(len( df[~(df['fare_amount'] > 1500)]))
```

```
1972319
1972319
```


```
[359]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='trip_distance', y='fare_amount',
↳ data=df[(df['trip_distance'] < 1)], alpha=0.5)
plt.title("Scatter Plot of Trip Distance vs Fare Amount")
plt.xlabel("Trip Distance (miles)")
plt.ylabel("Fare Amount ($)")
plt.grid(True)
plt.show()
```



```
[361]: # remove engtries with trip_distance is around 0.8 and fare_amount is more
        ↳ than 300
df = df[~((df['trip_distance'] < 0.8) & (df['fare_amount'] > 300))]
```

```
[367]: # remove engtries with trip_distance is 0.4 and fare_amount is more than 400
df = df[~((df['trip_distance'] < 0.4) & (df['fare_amount'] > 200))]
```

```
[371]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='trip_distance', y='fare_amount',
        ↳ data=df[(df['trip_distance'] > 1000)], alpha=0.5)
plt.title("Scatter Plot of Trip Distance vs Fare Amount")
plt.xlabel("Trip Distance (miles)")
plt.ylabel("Fare Amount ($)")
plt.grid(True)
plt.show()
```



```
[377]: # remove engtries with trip_distance motre then 10000 and fare_amount is less
↳ than 100
df = df[~((df['trip_distance'] > 10000) & (df['fare_amount'] < 200))]
```

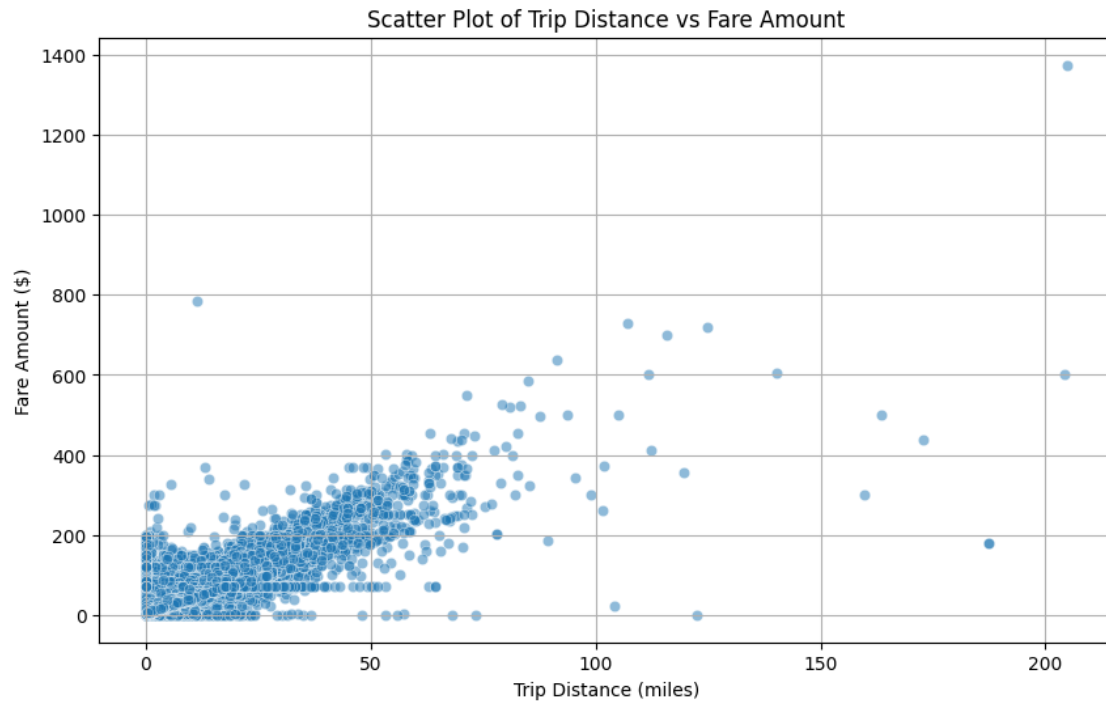
```
[393]: # remove engtries with trip_distance motre then 250
print(len(df[~((df['trip_distance'] > 250))]))
df = df[~((df['trip_distance'] > 250))]
```

1972185

```
[383]: df.shape
```

```
[383]: (1972205, 22)
```

```
[395]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='trip_distance', y='fare_amount', data=df, alpha=0.5)
plt.title("Scatter Plot of Trip Distance vs Fare Amount")
plt.xlabel("Trip Distance (miles)")
plt.ylabel("Fare Amount ($)")
plt.grid(True)
plt.show()
```

[400]: *# Do any columns need standardising?*

```
df.describe().loc[['min', 'max']]
df["payment_type"].value_counts()
# 0 is flex fare type
```

```
[400]: payment_type
1      1566923
2       332211
0       49225
4       14386
3        9440
Name: count, dtype: int64
```

1.7 3 Exploratory Data Analysis

[90 marks]

[402]: `df.columns.tolist()`

```
[402]: ['VendorID',
'tpep_pickup_datetime',
'tpep_dropoff_datetime',
'passenger_count',
'trip_distance',
```

```

'RatecodeID',
'store_and_fwd_flag',
'PULocationID',
'DOLocationID',
'payment_type',
'fare_amount',
'extra',
'mta_tax',
'tip_amount',
'tolls_amount',
'improvement_surcharge',
'total_amount',
'congestion_surcharge',
'airport_fee',
'pickup_date',
'pickup_hour',
'total_diff']

```

3.1 General EDA: Finding Patterns and Trends [40 marks]

3.1.1 [3 marks] Categorise the variables into Numerical or Categorical. * VendorID: Categorical * tpep_pickup_datetime: Numerical * tpep_dropoff_datetime: Numerical * passenger_count: Numerical * trip_distance: Numerical * RatecodeID: Categorical * PULocationID: Categorical * DOLocationID: Categorical * payment_type: Categorical * pickup_hour: Numerical * trip_duration: Numerical

The following monetary parameters belong in the same category, is it categorical or numerical?

Ans: All the below are numerical

- fare_amount
- extra
- mta_tax
- tip_amount
- tolls_amount
- improvement_surcharge
- total_amount
- congestion_surcharge
- airport_fee

Temporal Analysis 3.1.2 [5 marks] Analyse the distribution of taxi pickups by hours, days of the week, and months.

```

[412]: # Find and show the hourly trends in taxi pickups
# using line plot to look at the Pickup trends
hourly_trends = df['pickup_hour'].value_counts().sort_index()


plt.figure(figsize=(10, 6))

```

```

sns.lineplot(x=hourly_trends.index, y=hourly_trends.values, marker='o')
plt.title("Hourly Taxi Pickup Trends")
plt.xlabel("Hour of Day (0-23)")
plt.ylabel("Number of Pickups")
plt.xticks(range(0, 24))
plt.grid(True)
plt.show()

```



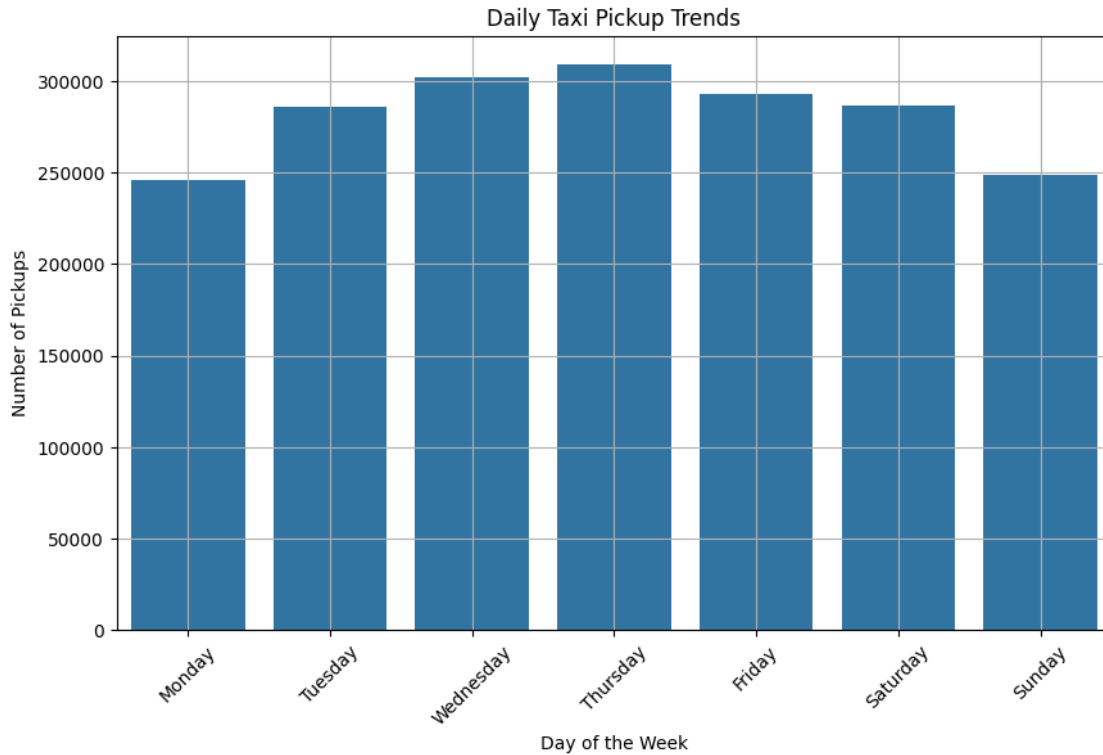
```

[416]: # Find and show the daily trends in taxi pickups (days of the week)
df['pickup_day'] = df['tpep_pickup_datetime'].dt.day_name()

daily_counts = df['pickup_day'].value_counts().reindex([
    'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday'
])

plt.figure(figsize=(10, 6))
sns.barplot(x=daily_counts.index, y=daily_counts.values)
plt.title("Daily Taxi Pickup Trends")
plt.xlabel("Day of the Week")
plt.ylabel("Number of Pickups")
plt.xticks(rotation=45)
plt.grid(True)
plt.show()


```



```
[418]: # Show the monthly trends in pickups

df['pickup_month'] = df['tpep_pickup_datetime'].dt.to_period('M').astype(str)
monthly_counts = df['pickup_month'].value_counts().sort_index()

plt.figure(figsize=(12, 6))
sns.barplot(x=monthly_counts.index, y=monthly_counts.values)
plt.title("Monthly Taxi Pickup Trends")
plt.xlabel("Month")
plt.ylabel("Number of Pickups")
plt.xticks(rotation=45)
plt.grid(True)
plt.show()
```



Financial Analysis Take a look at the financial parameters like `fare_amount`, `tip_amount`, `total_amount`, and also `trip_distance`. Do these contain zero/negative values?

```
[420]: # Analyse the above parameters

cols_to_check = ['fare_amount', 'tip_amount', 'total_amount', 'trip_distance']

for col in cols_to_check:
    zero_count = (df[col] == 0).sum()
    negative_count = (df[col] < 0).sum()
    print(f"{col}: {zero_count} zeros, {negative_count} negatives")
```

```
fare_amount: 605 zeros, 0 negatives
tip_amount: 451627 zeros, 0 negatives
total_amount: 270 zeros, 0 negatives
trip_distance: 35884 zeros, 0 negatives
```

Do you think it is beneficial to create a copy DataFrame leaving out the zero values from these?

3.1.3 [2 marks] Filter out the zero values from the above columns.

Note: The distance might be 0 in cases where pickup and drop is in the same zone. Do you think it is suitable to drop such cases of zero distance?

```
[425]: # Create a df with non zero entries for the selected parameters.
# Removing only those rows which have fare_amount or total amount as zero.

df = df[~((df['fare_amount'] == 0) | (df['total_amount'] == 0))]
```

```

for col in cols_to_check:
    zero_count = (df[col] == 0).sum()
    negative_count = (df[col] < 0).sum()
    print(f"{col}: {zero_count} zeros, {negative_count} negatives")

```

```

fare_amount: 0 zeros, 0 negatives
tip_amount: 451042 zeros, 0 negatives
total_amount: 0 zeros, 0 negatives
trip_distance: 35612 zeros, 0 negatives

```

```

[437]: zero_count = ((df['trip_distance'] == 0) & (df['fare_amount'] <= 25)).sum()
total_count = len(df)
percentage = (zero_count / total_count) * 100

print(f"Percentage of zero trip_distance records: {percentage:.2f}%")

```

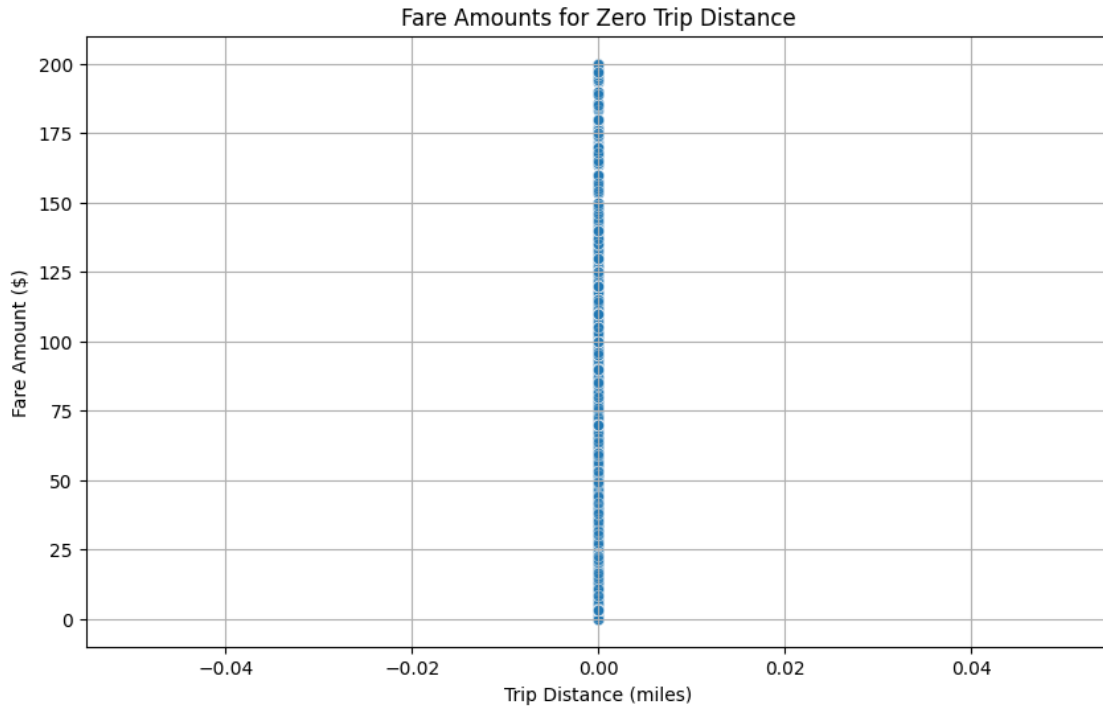
Percentage of zero trip_distance records: 1.23%

```

[427]: zero_distance_df = df[df['trip_distance'] == 0]

# Create scatter plot
plt.figure(figsize=(10, 6))
sns.scatterplot(x='trip_distance', y='fare_amount', data=zero_distance_df,
               alpha=0.6)
plt.title("Fare Amounts for Zero Trip Distance")
plt.xlabel("Trip Distance (miles)")
plt.ylabel("Fare Amount ($)")
plt.grid(True)
plt.show()

```



```
[439]: # only removing all the records with trip distance as 0 and amount less 25
        ↪ dollars
df = df[~ ((df['trip_distance'] == 0) & (df['fare_amount'] <= 25 ))]
```

```
[441]: zero_count = ((df['trip_distance'] == 0) ).sum()
total_count = len(df)
percentage = (zero_count / total_count) * 100

print(f"Percentage of zero trip_distance records: {percentage:.2f}%")
```

Percentage of zero trip_distance records: 0.58%


3.1.4 [3 marks] Analyse the monthly revenue (total_amount) trend

```
[449]: # Group data by month and analyse monthly revenue
import matplotlib.ticker as ticker

monthly_revenue = df.groupby('pickup_month')['total_amount'].sum()
plt.figure(figsize=(12, 6))
sns.lineplot(x=monthly_revenue.index, y=monthly_revenue.values, marker='o')
plt.gca().yaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: f'{x/
        ↪1000:.0f}K'))

plt.title("Monthly Taxi Revenue")
plt.xlabel("Month")
```

```
plt.ylabel("Total Revenue ($)")
plt.xticks(rotation=45)
plt.grid(True)
plt.show()
```



3.1.5 [3 marks] Show the proportion of each quarter of the year in the revenue

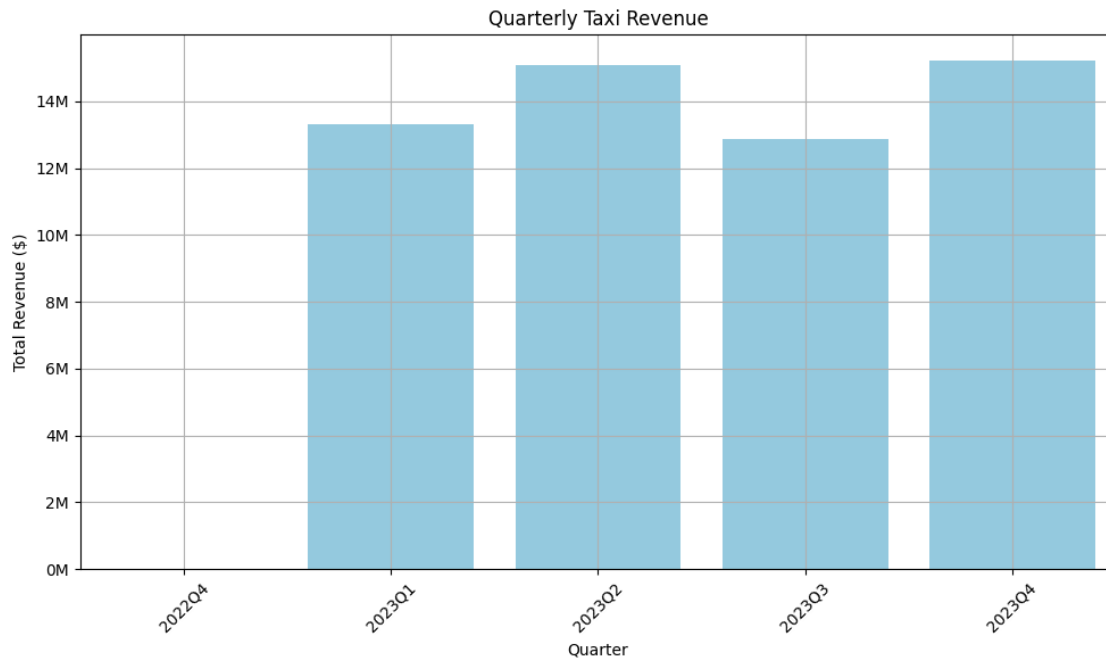
```
[451]: # Calculate proportion of each quarter
df['pickup_quarter'] = df['tpep_pickup_datetime'].dt.to_period('Q')
quarterly_revenue = df.groupby('pickup_quarter')['total_amount'].sum()
annual_revenue = quarterly_revenue.sum()
quarterly_proportion = (quarterly_revenue / annual_revenue) * 100
print(quarterly_proportion.round(2))
```

```
pickup_quarter
2022Q4      0.00
2023Q1     23.56
2023Q2     26.72
2023Q3     22.76
2023Q4     26.96
Freq: Q-DEC, Name: total_amount, dtype: float64
```

```
[457]: plt.figure(figsize=(10, 6))
sns.barplot(x=quarterly_revenue.index.astype(str), y=quarterly_revenue.values,
            color='skyblue')
plt.gca().yaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: f'{x/
            1000000:.0f}M'))
```




```
plt.title("Quarterly Taxi Revenue")
plt.xlabel("Quarter")
plt.ylabel("Total Revenue ($)")
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```



3.1.6 [3 marks] Visualise the relationship between `trip_distance` and `fare_amount`. Also find the correlation value for these two.

Hint: You can leave out the trips with `trip_distance = 0`

```
[463]: # Show how trip fare is affected by distance
plt.figure(figsize=(10, 6))
sns.scatterplot(x='trip_distance', y='fare_amount', data=df[(df['trip_distance'] > 25)], alpha=0.4)
plt.title("Trip Distance vs Fare Amount")
plt.xlabel("Trip Distance (miles)")
plt.ylabel("Fare Amount ($)")
plt.grid(True)
plt.show()
```



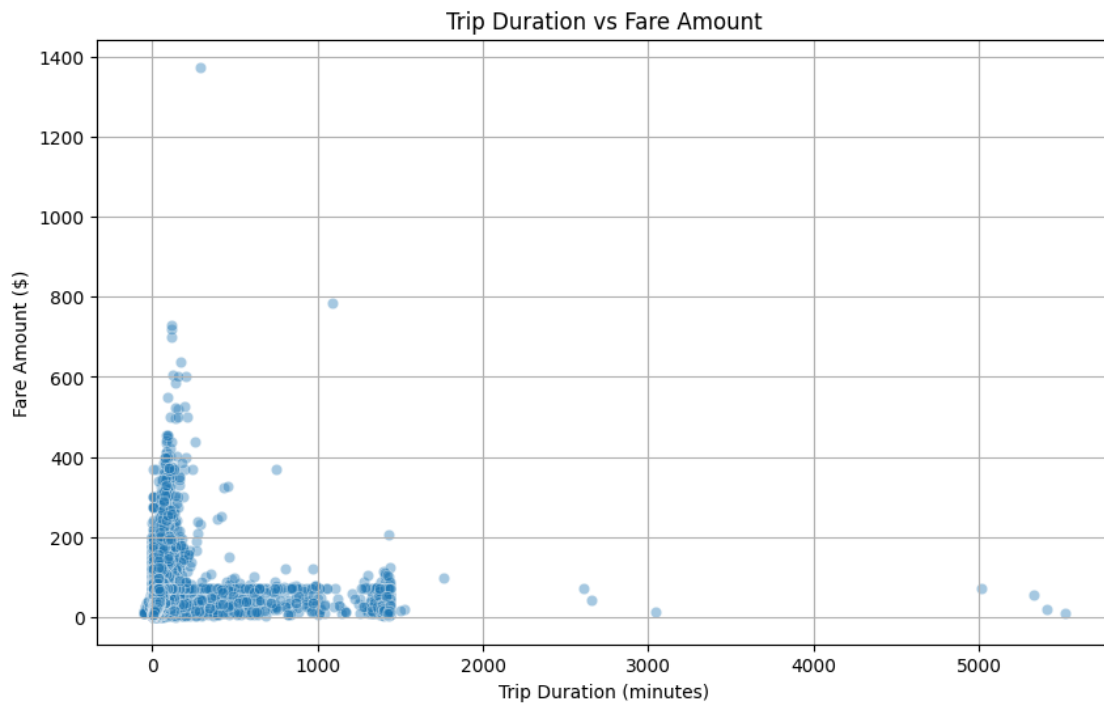
3.1.7 [5 marks] Find and visualise the correlation between: 1. `fare_amount` and trip duration (pickup time to dropoff time) 2. `fare_amount` and `passenger_count` 3. `tip_amount` and `trip_distance`

```
[467]: df.columns
```

```
[467]: Index(['VendorID', 'tpep_pickup_datetime', 'tpep_dropoff_datetime',
        'passenger_count', 'trip_distance', 'RatecodeID', 'store_and_fwd_flag',
        'PULocationID', 'DOLocationID', 'payment_type', 'fare_amount', 'extra',
        'mta_tax', 'tip_amount', 'tolls_amount', 'improvement_surcharge',
        'total_amount', 'congestion_surcharge', 'airport_fee', 'pickup_date',
        'pickup_hour', 'total_diff', 'pickup_day', 'pickup_month',
        'pickup_quarter'],
        dtype='object')
```


```
[469]: # Show relationship between fare and trip duration
df['trip_duration'] = (df['tpep_dropoff_datetime'] -
    df['tpep_pickup_datetime']).dt.total_seconds() / 60
plt.figure(figsize=(10, 6))
sns.scatterplot(x='trip_duration', y='fare_amount', data=df, alpha=0.4)
plt.title("Trip Duration vs Fare Amount")
plt.xlabel("Trip Duration (minutes)")
plt.ylabel("Fare Amount ($)")
plt.grid(True)
```

```
plt.show()
```




```
[475]: # Show relationship between fare and number of passengers
```

```
plt.figure(figsize=(10, 6))
sns.boxplot(x='passenger_count', y='fare_amount', data=df)
plt.title("Fare Amount by Number of Passengers")
plt.xlabel("Passenger Count")
plt.ylabel("Fare Amount ($)")
plt.grid(True)
plt.show()
```



```
[477]: # Show relationship between tip and trip distance
```

```
plt.figure(figsize=(10, 6))
sns.scatterplot(x='trip_distance', y='tip_amount', data=df, alpha=0.4)
plt.title("Tip Amount vs Trip Distance")
plt.xlabel("Trip Distance (miles)")
plt.ylabel("Tip Amount ($)")
plt.grid(True)
plt.show()
```



3.1.8 [3 marks] Analyse the distribution of different payment types (`payment_type`)

```
[485]: payment_labels = {
        0: 'Flex Fare',
        1: 'Credit Card',
        2: 'Cash',
        3: 'No Charge',
        4: 'Dispute'
    }
df['payment_type_label'] = df['payment_type'].map(payment_labels)

[487]: # Analyse the distribution of different payment types (payment_type).


payment_counts = df['payment_type_label'].value_counts().sort_index()

plt.figure(figsize=(8, 5))
sns.barplot(x=payment_counts.index, y=payment_counts.values, palette='pastel')
plt.title("Distribution of Payment Types")
plt.xlabel("Payment Method")
plt.ylabel("Number of Transactions")
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```

C:\Users\Sherry\AppData\Local\Temp\ipykernel_11388\3317766951.py:6:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=payment_counts.index, y=payment_counts.values, palette='pastel')
```




```
[491]: explode = [0.05] * len(payment_counts) # Adjust the value for more/less separation

# Plot pie chart
plt.figure(figsize=(6, 6))
plt.pie(payment_counts.values,
        labels=payment_counts.index,
        autopct='%1.1f%%',
        startangle=140,
        explode=explode,
        colors=sns.color_palette('pastel'))

plt.title("Payment Type Distribution (Exploded View)")
plt.tight_layout()
plt.show()
```

Payment Type Distribution (Exploded View)



- 1= Credit card
- 2= Cash
- 3= No charge
- 4= Dispute

Geographical Analysis For this, you have to use the *taxi_zones.shp* file from the *taxi_zones* folder.

There would be multiple files inside the folder (such as *.shx*, *.sbx*, *.sbn* etc). You do not need to import/read any of the files other than the shapefile, *taxi_zones.shp*.

Do not change any folder structure - all the files need to be present inside the folder for it to work.

The folder structure should look like this:

```
Taxi Zones
|- taxi_zones.shp.xml
```

```
|- taxi_zones.prj
|- taxi_zones.sbn
|- taxi_zones.shp
|- taxi_zones.dbf
|- taxi_zones.shx
|- taxi_zones.sbx
```

You only need to read the `taxi_zones.shp` file. The `shp` file will utilise the other files by itself.

We will use the *GeoPandas* library for geopgraphical analysis

```
import geopandas as gpd
```

More about geopandas and shapefiles: [About](#)

Reading the shapefile is very similar to *Pandas*. Use `gpd.read_file()` function to load the data (`taxi_zones.shp`) as a `GeoDataFrame`. Documentation: [Reading and Writing Files](#)

```
[493]: !pip install geopandas
```

Collecting geopandas

Downloading geopandas-1.1.0-py3-none-any.whl.metadata (2.3 kB)

Requirement already satisfied: numpy>=1.24 in

c:\users\sherry\anaconda3\lib\site-packages (from geopandas) (1.26.4)

Collecting pyogrio>=0.7.2 (from geopandas)

Downloading pyogrio-0.11.0-cp312-cp312-win_amd64.whl.metadata (5.4 kB)

Requirement already satisfied: packaging in

c:\users\sherry\appdata\roaming\python\python312\site-packages (from geopandas) (24.2)

Requirement already satisfied: pandas>=2.0.0 in

c:\users\sherry\appdata\roaming\python\python312\site-packages (from geopandas) (2.2.3)

Collecting pyproj>=3.5.0 (from geopandas)

Downloading pyproj-3.7.1-cp312-cp312-win_amd64.whl.metadata (31 kB)

Collecting shapely>=2.0.0 (from geopandas)

Downloading shapely-2.1.1-cp312-cp312-win_amd64.whl.metadata (7.0 kB)

Requirement already satisfied: python-dateutil>=2.8.2 in

c:\users\sherry\appdata\roaming\python\python312\site-packages (from pandas>=2.0.0->geopandas) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in

c:\users\sherry\appdata\roaming\python\python312\site-packages (from pandas>=2.0.0->geopandas) (2024.2)

Requirement already satisfied: tzdata>=2022.7 in

c:\users\sherry\appdata\roaming\python\python312\site-packages (from pandas>=2.0.0->geopandas) (2024.2)

Requirement already satisfied: certifi in

c:\users\sherry\appdata\roaming\python\python312\site-packages (from pyogrio>=0.7.2->geopandas) (2025.1.31)

Requirement already satisfied: six>=1.5 in

c:\users\sherry\appdata\roaming\python\python312\site-packages (from python-dateutil>=2.8.2->pandas>=2.0.0->geopandas) (1.17.0)


```

Downloading geopandas-1.1.0-py3-none-any.whl (338 kB)
Downloading pyogrio-0.11.0-cp312-cp312-win_amd64.whl (19.2 MB)
----- 0.0/19.2 MB ? eta -:--:--
----- 4.2/19.2 MB 22.9 MB/s eta 0:00:01
----- 9.2/19.2 MB 23.8 MB/s eta 0:00:01
----- 12.6/19.2 MB 20.7 MB/s eta 0:00:01
----- 17.0/19.2 MB 21.1 MB/s eta 0:00:01
----- 19.2/19.2 MB 20.9 MB/s eta 0:00:00
Downloading pyproj-3.7.1-cp312-cp312-win_amd64.whl (6.3 MB)
----- 0.0/6.3 MB ? eta -:--:--
----- 4.5/6.3 MB 24.4 MB/s eta 0:00:01
----- 6.3/6.3 MB 19.2 MB/s eta 0:00:00
Downloading shapely-2.1.1-cp312-cp312-win_amd64.whl (1.7 MB)
----- 0.0/1.7 MB ? eta -:--:--
----- 1.7/1.7 MB 22.9 MB/s eta 0:00:00
Installing collected packages: shapely, pyproj, pyogrio, geopandas
Successfully installed geopandas-1.1.0 pyogrio-0.11.0 pyproj-3.7.1 shapely-2.1.1

```

[notice] A new release of pip is available: 24.3.1 -> 25.1.1
[notice] To update, run: python.exe -m pip install --upgrade pip

3.1.9 [2 marks] Load the shapefile and display it.

```

[499]: import geopandas as gpd

# Read the shapefile using geopandas
zones = gpd.read_file(r"C:\Users\Sherry\UpgradExample\Datasets and_
↳Dictionary-NYC\Starter Notebook - EDA NYC Taxi\taxi_zones\taxi_zones.shp")
# read the .shp file using gpd
zones.head()

```

```

[499]:
  OBJECTID  Shape_Leng  Shape_Area      zone  LocationID  \
0         1    0.116357    0.000782  Newark Airport      1
1         2    0.433470    0.004866   Jamaica Bay      2
2         3    0.084341    0.000314 Allerton/Pelham Gardens  3
3         4    0.043567    0.000112   Alphabet City      4
4         5    0.092146    0.000498   Arden Heights      5

```

```

      borough      geometry
0      EWR  POLYGON ((933100.918 192536.086, 933091.011 19...
1    Queens  MULTIPOLYGON (((1033269.244 172126.008, 103343...
2    Bronx  POLYGON ((1026308.77 256767.698, 1026495.593 2...
3  Manhattan  POLYGON ((992073.467 203714.076, 992068.667 20...
4  Staten Island  POLYGON ((935843.31 144283.336, 936046.565 144...

```

Now, if you look at the DataFrame created, you will see columns like: OBJECTID, Shape_Leng, Shape_Area, zone, LocationID, borough, geometry.

Now, the `locationID` here is also what we are using to mark pickup and drop zones in the trip records.


The geometric parameters like shape length, shape area and geometry are used to plot the zones on a map.

This can be easily done using the `plot()` method.

```
[501]: print(zones.info())
       zones.plot()
```

```
<class 'geopandas.geodataframe.GeoDataFrame'>
RangeIndex: 263 entries, 0 to 262
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   OBJECTID    263 non-null   int32
1   Shape_Leng  263 non-null   float64
2   Shape_Area  263 non-null   float64
3   zone        263 non-null   object
4   LocationID  263 non-null   int32
5   borough     263 non-null   object
6   geometry    263 non-null   geometry
dtypes: float64(2), geometry(1), int32(2), object(2)
memory usage: 12.5+ KB
None
```

```
[501]: <Axes: >
```



Now, you have to merge the trip records and zones data using the location IDs.

3.1.10 [3 marks] Merge the zones data into trip data using the `locationID` and `PULocationID` columns.

```
[503]: # Merge zones and trip records using locationID and PULocationID

df = df.merge(zones[['LocationID', 'zone', 'borough']], how='left',
              ↪left_on='PULocationID', right_on='LocationID')
```

```
[507]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1947483 entries, 0 to 1947482
Data columns (total 30 columns):
#   Column                Dtype
---  -
0   VendorID              int64
1   tpep_pickup_datetime  datetime64[us]
2   tpep_dropoff_datetime datetime64[us]
3   passenger_count       float64
4   trip_distance         float64
```

```

5   RatecodeID           float64
6   store_and_fwd_flag   object
7   PULocationID         int64
8   DOLocationID         int64
9   payment_type         int64
10  fare_amount          float64
11  extra                float64
12  mta_tax              float64
13  tip_amount           float64
14  tolls_amount         float64
15  improvement_surcharge float64
16  total_amount         float64
17  congestion_surcharge float64
18  airport_fee          float64
19  pickup_date          object
20  pickup_hour          int32
21  total_diff           float64
22  pickup_day           object
23  pickup_month         object
24  pickup_quarter       period[Q-DEC]
25  trip_duration        float64
26  payment_type_label   object
27  LocationID           float64
28  zone                 object
29  borough              object
dtypes: datetime64[us](2), float64(15), int32(1), int64(4), object(7),
period[Q-DEC](1)
memory usage: 438.3+ MB

```

3.1.11 [3 marks] Group data by location IDs to find the total number of trips per location ID

```

[511]: # Group data by location and calculate the number of trips

pickup = df['PULocationID'].value_counts().rename_axis('LocationID').
        ↪reset_index(name='pickup_count')
dropoff = df['DOLocationID'].value_counts().rename_axis('LocationID').
        ↪reset_index(name='dropoff_count')

# Merge and fill missing values with 0
location_trip_counts = pd.merge(pickup, dropoff, on='LocationID', how='outer').
        ↪fillna(0)

# Calculate total trips
location_trip_counts['total_trips'] = location_trip_counts['pickup_count'] +
        ↪location_trip_counts['dropoff_count']
location_trip_counts

```

```
[511]:
```

	LocationID	pickup_count	dropoff_count	total_trips
0	1	198.0	5593.0	5791.0
1	2	2.0	4.0	6.0
2	3	39.0	164.0	203.0
3	4	2273.0	7370.0	9643.0
4	5	13.0	34.0	47.0
..
256	261	10289.0	9280.0	19569.0
257	262	25874.0	30149.0	56023.0
258	263	37490.0	40052.0	77542.0
259	264	17096.0	18275.0	35371.0
260	265	890.0	8016.0	8906.0

[261 rows x 4 columns]

3.1.12 [2 marks] Now, use the grouped data to add number of trips to the GeoDataFrame.

We will use this to plot a map of zones showing total trips per zone.

```
[514]: # Merge trip counts back to the zones GeoDataFrame

zones = zones.merge(location_trip_counts, on='LocationID', how='left')
zones.head()
```

```
[514]:
```

	OBJECTID	Shape_Leng	Shape_Area	zone	LocationID	\
0	1	0.116357	0.000782	Newark Airport	1	
1	2	0.433470	0.004866	Jamaica Bay	2	
2	3	0.084341	0.000314	Allerton/Pelham Gardens	3	
3	4	0.043567	0.000112	Alphabet City	4	
4	5	0.092146	0.000498	Arden Heights	5	

	borough	geometry	\
0	EWB	POLYGON ((933100.918 192536.086, 933091.011 19...	
1	Queens	MULTIPOLYGON (((1033269.244 172126.008, 103343...	
2	Bronx	POLYGON ((1026308.77 256767.698, 1026495.593 2...	
3	Manhattan	POLYGON ((992073.467 203714.076, 992068.667 20...	
4	Staten Island	POLYGON ((935843.31 144283.336, 936046.565 144...	

	pickup_count	dropoff_count	total_trips
0	198.0	5593.0	5791.0
1	2.0	4.0	6.0
2	39.0	164.0	203.0
3	2273.0	7370.0	9643.0
4	13.0	34.0	47.0

The next step is creating a color map (choropleth map) showing zones by the number of trips taken.

Again, you can use the `zones.plot()` method for this. [Plot Method GPD](#)

But first, you need to define the figure and axis for the plot.

```
fig, ax = plt.subplots(1, 1, figsize = (12, 10))
```

This function creates a figure (fig) and a single subplot (ax)

After setting up the figure and axis, we can proceed to plot the GeoDataFrame on this axis. This is done in the next step where we use the plot method of the GeoDataFrame.

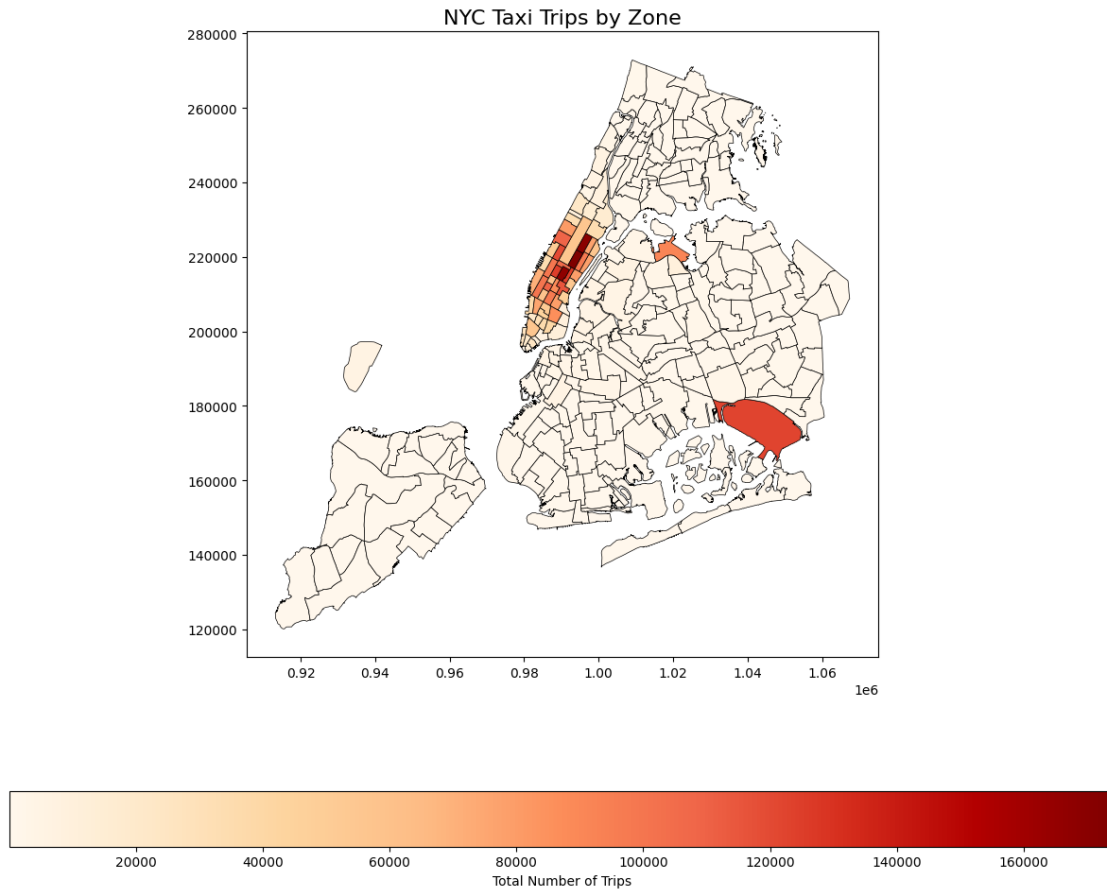
You can define the following parameters in the `zones.plot()` method:

```
column = '',  
ax = ax,  
legend = True,  
legend_kwds = {'label': "label", 'orientation': "<horizontal/vertical>"}
```

To display the plot, use `plt.show()`.

3.1.13 [3 marks] Plot a color-coded map showing zone-wise trips

```
[520]: # Define figure and axis  
fig, ax = plt.subplots(1, 1, figsize=(12, 10))  
  
# Plot the map and display it  
zones.plot(column='total_trips',  
           cmap='OrRd',  
           legend=True,  
           edgecolor='black',  
           linewidth=0.5,  
           ax=ax,  
           legend_kwds={  
               'label': "Total Number of Trips",  
               'orientation': "horizontal" # or "vertical"  
           })  
  
ax.set_title("NYC Taxi Trips by Zone", fontsize=16)  
plt.tight_layout()  
plt.show()
```



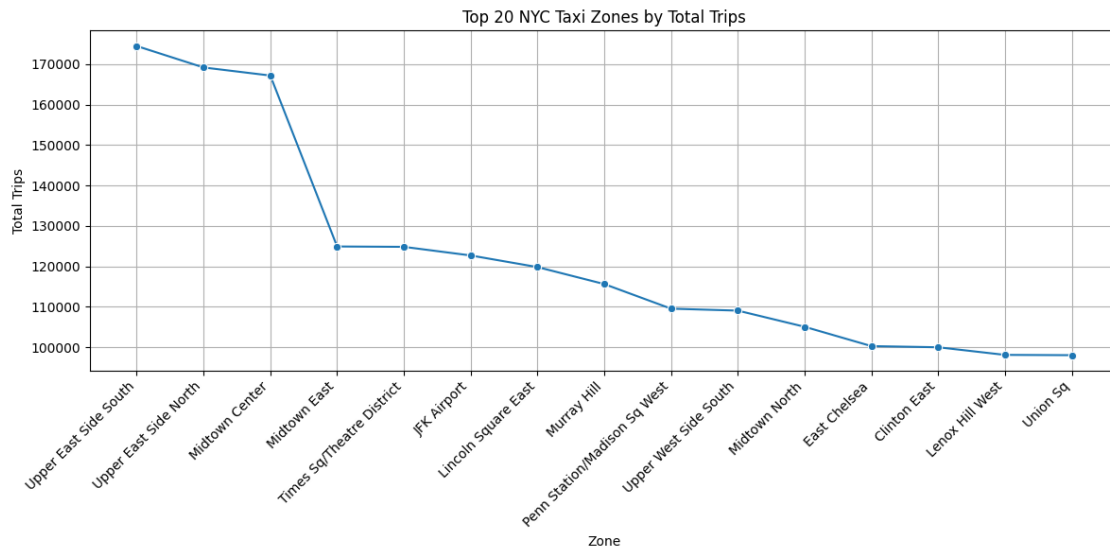
[524]: *# can you try displaying the zones DF sorted by the number of trips?*

```
sorted_zones = zones.sort_values(by='total_trips', ascending=False)

# Plot top 20 zones (or adjust as needed)
plt.figure(figsize=(12, 6))
sns.lineplot(
    x=sorted_zones['zone'].head(15),
    y=sorted_zones['total_trips'].head(15),
    marker='o'
)

plt.title("Top 20 NYC Taxi Zones by Total Trips")
plt.xlabel("Zone")
plt.ylabel("Total Trips")
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.grid(True)
```

```
plt.show()
```



Here we have completed the temporal, financial and geographical analysis on the trip records.

Compile your findings from general analysis below:

You can consider the following points:

- Busiest hours, days and months
- Trends in revenue collected
- Trends in quarterly revenue
- How fare depends on trip distance, trip duration and passenger counts
- How tip amount depends on trip distance
- Busiest zones

3.2 Detailed EDA: Insights and Strategies [50 marks]

Having performed basic analyses for finding trends and patterns, we will now move on to some detailed analysis focussed on operational efficiency, pricing strategies, and customer experience.

Operational Efficiency Analyze variations by time of day and location to identify bottlenecks or inefficiencies in routes

3.2.1 [3 marks] Identify slow routes by calculating the average time taken by cabs to get from one zone to another at different hours of the day.

Speed on a route X for hour $Y = (\text{distance of the route } X / \text{average trip duration for hour } Y)$

```
[526]: # Find routes which have the slowest speeds at different times of the day

df['trip_duration_hr'] = df['trip_duration'] / 60 # convert minutes to hours
df['speed_mph'] = df['trip_distance'] / df['trip_duration_hr']
```



```
[540]: negative_speed_rows = df[df['speed_mph'] < 0]
len(negative_speed_rows)
#remove negative speed rows from the data set
df = df[df['speed_mph'] >= 0]
```

```
[528]: def get_time_of_day(hour):
    if 5 <= hour < 12:
        return 'Morning'
    elif 12 <= hour < 17:
        return 'Afternoon'
    elif 17 <= hour < 21:
        return 'Evening'
    else:
        return 'Night'

df['time_of_day'] = df['pickup_hour'].apply(get_time_of_day)
df['route'] = df['zone'] + " " + df['DOLocationID'].astype(str)
```

```
[528]:
```

		route	time_of_day	speed_mph
0	Allerton/Pelham Gardens	10	Morning	17.387074
1	Allerton/Pelham Gardens	119	Evening	0.000000
2	Allerton/Pelham Gardens	133	Morning	0.000000
3	Allerton/Pelham Gardens	137	Morning	14.452188
4	Allerton/Pelham Gardens	142	Night	12.731092
...	
44716	Yorkville West	95	Night	33.024821
44717	Yorkville West	97	Afternoon	16.922724
44718	Yorkville West	97	Evening	18.859933
44719	Yorkville West	97	Morning	21.771459
44720	Yorkville West	97	Night	23.045805

[44721 rows x 3 columns]

```
[542]: route_speeds = df[(df['speed_mph'] > 0)].groupby(['route', 'time_of_day'])['speed_mph'].mean().reset_index()
route_speeds
```

```
[542]:
```

		route	time_of_day	speed_mph
0	Allerton/Pelham Gardens	10	Morning	17.387074
1	Allerton/Pelham Gardens	137	Morning	14.452188
2	Allerton/Pelham Gardens	142	Night	12.731092
3	Allerton/Pelham Gardens	147	Afternoon	14.716981
4	Allerton/Pelham Gardens	163	Morning	17.384074
...	
43366	Yorkville West	95	Night	33.024821
43367	Yorkville West	97	Afternoon	16.922724
43368	Yorkville West	97	Evening	18.859933

43369	Yorkville West	97	Morning	21.771459
43370	Yorkville West	97	Night	23.045805

[43371 rows x 3 columns]

```
[546]: #Below are the slowest routes at different time of day
slowest_routes = route_speeds.sort_values(['time_of_day', 'speed_mph']).
↳groupby('time_of_day').head(1)
print(slowest_routes)
```

		route	time_of_day	speed_mph
16515	Howard Beach	129	Afternoon	0.072029
39723	Washington Heights North	264	Evening	0.007772
4109	Central Park	10	Morning	0.022236
15307	Greenwich Village North	235	Night	0.048105

How does identifying high-traffic, high-demand routes help us?

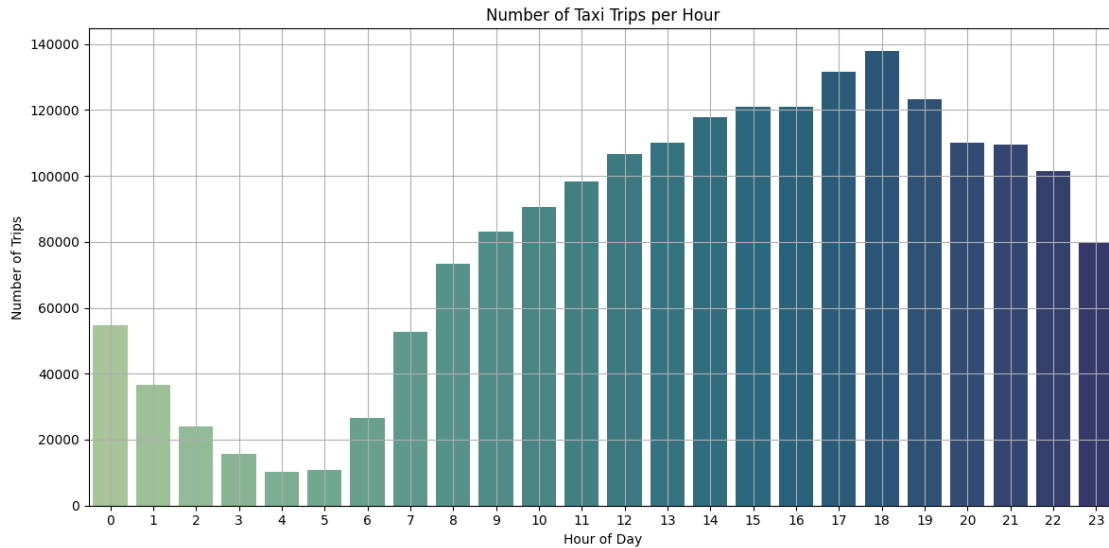
3.2.2 [3 marks] Calculate the number of trips at each hour of the day and visualise them. Find the busiest hour and show the number of trips for that hour.

```
[561]: # Visualise the number of trips per hour and find the busiest hour

hourly_counts = df['pickup_hour'].value_counts().sort_index()

# Plot
plt.figure(figsize=(12, 6))
sns.barplot(x=hourly_counts.index, y=hourly_counts.values, hue=hourly_counts.
↳index, palette='crest', legend=False)

plt.title("Number of Taxi Trips per Hour")
plt.xlabel("Hour of Day")
plt.ylabel("Number of Trips")
plt.xticks(range(0, 24))
plt.grid(True)
plt.tight_layout()
plt.show()
```



Busiest time is between 6 - 7 PM.

Remember, we took a fraction of trips. To find the actual number, you have to scale the number up by the sampling ratio.

3.2.3 [2 mark] Find the actual number of trips in the five busiest hours

```
[565]: hourly_counts = df['pickup_hour'].value_counts().sort_values(ascending=False)
top_5_hours = hourly_counts.head(5)

# Scale up the number of trips

# Fill in the value of your sampling fraction and use that to scale up the
↳ numbers
sample_fraction = 0.05
estimated_top_5 = (top_5_hours / sample_fraction).astype(int)

print("Estimated Total Trips in Top 5 Busiest Hours:")
print(estimated_top_5)
```

Estimated Total Trips in Top 5 Busiest Hours:

```
pickup_hour
18    2755620
17    2629080
19    2466940
16    2420800
15    2420100
Name: count, dtype: int32
```

3.2.4 [3 marks] Compare hourly traffic pattern on weekdays. Also compare for weekend.

```
[568]: # Compare traffic trends for the week days and weekends
df['pickup_dayofweek'] = df['tpep_pickup_datetime'].dt.dayofweek # Monday=0,
↳ Sunday=6
df['day_type'] = df['pickup_dayofweek'].apply(lambda x: 'Weekend' if x >= 5
↳ else 'Weekday')
```

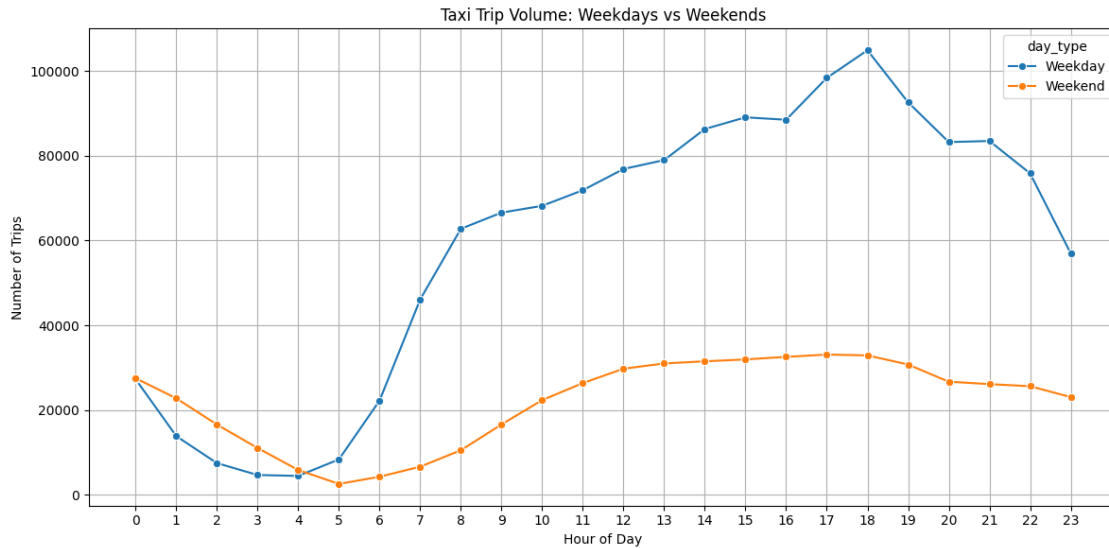
```
[570]: hourly_trends = df.groupby(['pickup_hour', 'day_type']).size().
↳ reset_index(name='trip_count')
hourly_trends.head(10)
```

```
[570]:
```

	pickup_hour	day_type	trip_count
0	0	Weekday	27286
1	0	Weekend	27518
2	1	Weekday	13865
3	1	Weekend	22779
4	2	Weekday	7496
5	2	Weekend	16592
6	3	Weekday	4694
7	3	Weekend	11059
8	4	Weekday	4478
9	4	Weekend	5874

```
[572]: plt.figure(figsize=(12, 6))
sns.lineplot(data=hourly_trends, x='pickup_hour', y='trip_count',
↳ hue='day_type', marker='o')

plt.title("Taxi Trip Volume: Weekdays vs Weekends")
plt.xlabel("Hour of Day")
plt.ylabel("Number of Trips")
plt.xticks(range(0, 24))
plt.grid(True)
plt.tight_layout()
plt.show()
```



Trends on weekday vs Weekend There is more traffic at late night 12 AM during weekends vs the weekadys. In general between 4 to 11 PM traffic is high on weekdays may be because of business or employees travelling to there home.

What can you infer from the above patterns? How will finding busy and quiet hours for each day help us?

3.2.5 [3 marks] Identify top 10 zones with high hourly pickups. Do the same for hourly dropoffs. Show pickup and dropoff trends in these zones.

```
[579]: # Find top 10 pickup and dropoff zones

pickup_hourly = df.groupby(['zone', 'pickup_hour']).size().
    ↳reset_index(name='pickup_count')
dropoff_hourly = df.groupby(['DOLocationID', 'pickup_hour']).size().
    ↳reset_index(name='dropoff_count')
dropoff_hourly = dropoff_hourly.merge(zones[['LocationID', 'zone']],
    ↳left_on='DOLocationID', right_on='LocationID', how='left')
dropoff_hourly.drop(columns='LocationID', inplace=True)

top_pickup_zones = pickup_hourly.groupby('zone')['pickup_count'].sum().
    ↳nlargest(10).index
print (top_pickup_zones)
top_dropoff_zones = dropoff_hourly.groupby('zone')['dropoff_count'].sum().
    ↳nlargest(10).index
print (top_dropoff_zones)
```

```
Index(['JFK Airport', 'Upper East Side South', 'Midtown Center',
      'Upper East Side North', 'Midtown East', 'LaGuardia Airport',
      'Penn Station/Madison Sq West', 'Times Sq/Theatre District',
```

```

        'Lincoln Square East', 'Murray Hill'],
        dtype='object', name='zone')
Index(['Upper East Side North', 'Upper East Side South', 'Midtown Center',
       'Times Sq/Theatre District', 'Murray Hill', 'Midtown East',
       'Lincoln Square East', 'Upper West Side South', 'Lenox Hill West',
       'East Chelsea'],
      dtype='object', name='zone')

```

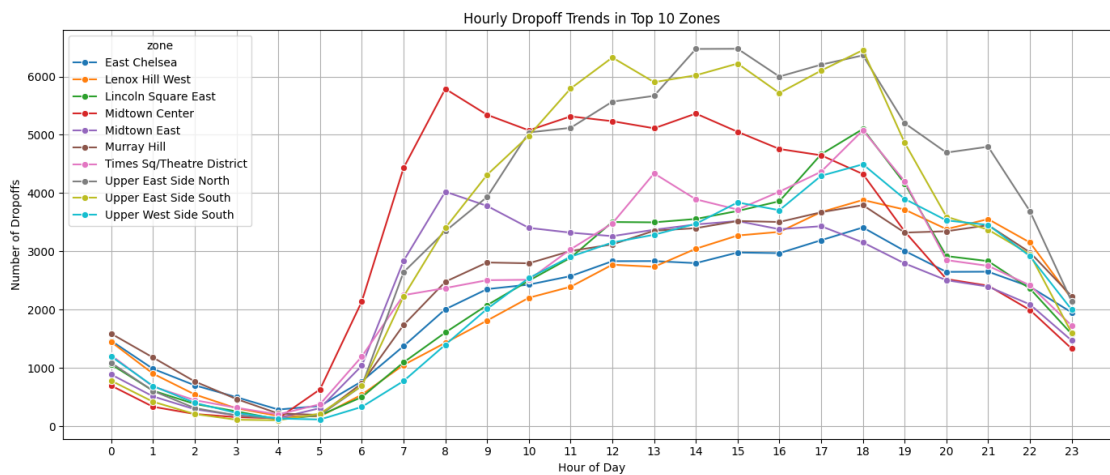
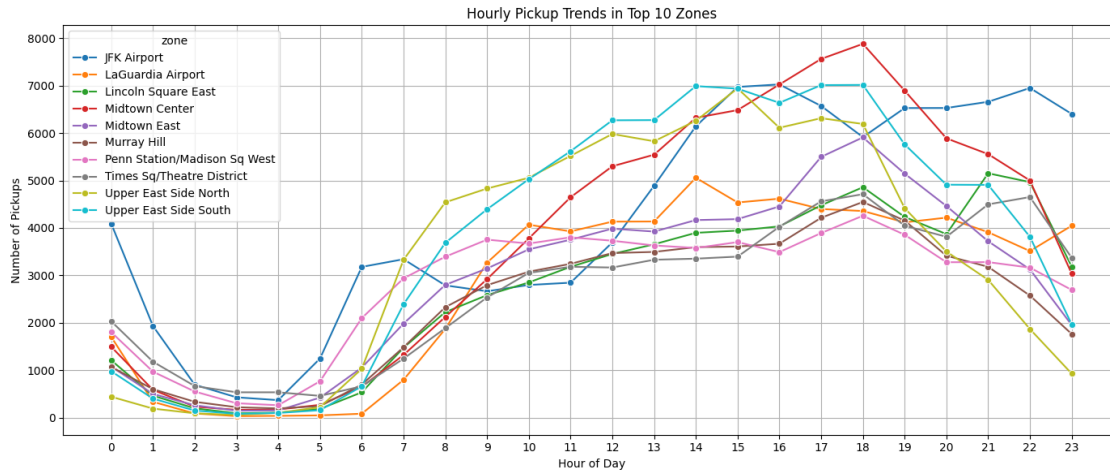
```

[582]: pickup_top = pickup_hourly[pickup_hourly['zone'].isin(top_pickup_zones)]
       dropoff_top = dropoff_hourly[dropoff_hourly['zone'].isin(top_dropoff_zones)]

       # Plot pickup trends
       plt.figure(figsize=(14, 6))
       sns.lineplot(data=pickup_top, x='pickup_hour', y='pickup_count', hue='zone',
                    ↪marker='o')
       plt.title("Hourly Pickup Trends in Top 10 Zones")
       plt.xlabel("Hour of Day")
       plt.ylabel("Number of Pickups")
       plt.xticks(range(0, 24))
       plt.grid(True)
       plt.tight_layout()
       plt.show()

       # Plot dropoff trends
       plt.figure(figsize=(14, 6))
       sns.lineplot(data=dropoff_top, x='pickup_hour', y='dropoff_count', hue='zone',
                    ↪marker='o')
       plt.title("Hourly Dropoff Trends in Top 10 Zones")
       plt.xlabel("Hour of Day")
       plt.ylabel("Number of Dropoffs")
       plt.xticks(range(0, 24))
       plt.grid(True)
       plt.tight_layout()
       plt.show()

```



3.2.6 [3 marks] Find the ratio of pickups and dropoffs in each zone. Display the 10 highest (pickup/drop) and 10 lowest (pickup/drop) ratios.

[640]: *# Find the top 10 and bottom 10 pickup/dropoff ratios*

```
pickup_counts = df.groupby('PULocationID').size().
    ↪reset_index(name='pickup_count')
pickup_counts.rename(columns={'PULocationID': 'LocationID'}, inplace=True)
# Count dropoffs
dropoff_counts = df.groupby('DOLocationID').size().
    ↪reset_index(name='dropoff_count')
dropoff_counts.rename(columns={'DOLocationID': 'LocationID'}, inplace=True)
# Merge counts
zone_counts = pd.merge(pickup_counts, dropoff_counts, on='LocationID',
    ↪how='outer').fillna(0)
```

```

# Avoid division by zero
zone_counts = zone_counts[(zone_counts['dropoff_count'] > 0) &
    ↪(zone_counts['pickup_count'] > 0)]

# Compute pickup/dropoff ratio
zone_counts['pickup_dropoff_ratio'] = zone_counts['pickup_count'] /
    ↪zone_counts['dropoff_count']

#Add zone column
zone_counts = zone_counts.merge(zones[['LocationID', 'zone']], on='LocationID',
    ↪how='left')

# Top 10 highest ratios
top10 = zone_counts.sort_values(by='pickup_dropoff_ratio', ascending=False).
    ↪head(10)

# Bottom 10 lowest ratios
bottom10 = zone_counts.sort_values(by='pickup_dropoff_ratio', ascending=True).
    ↪head(10)

print("Top 10 Pickup/Dropoff Ratios:")
print(top10[['zone', 'pickup_count', 'dropoff_count', 'pickup_dropoff_ratio']])

print("\n Bottom 10 Pickup/Dropoff Ratios:")
print(bottom10[['zone', 'pickup_count', 'dropoff_count',
    ↪'pickup_dropoff_ratio']])

```

Top 10 Pickup/Dropoff Ratios:

	zone	pickup_count	dropoff_count	\
69	East Elmhurst	8681.0	943.0	
126	JFK Airport	100661.0	21981.0	
132	LaGuardia Airport	67284.0	24361.0	
179	Penn Station/Madison Sq West	66844.0	42673.0	
108	Greenwich Village South	25903.0	18780.0	
41	Central Park	32558.0	23716.0	
239	West Village	43322.0	32497.0	
156	Midtown East	69352.0	55550.0	
155	Midtown Center	90851.0	76350.0	
98	Garment District	31913.0	26822.0	

	pickup_dropoff_ratio
69	9.205726
126	4.579455
132	2.761956
179	1.566424
108	1.379286

41	1.372828
239	1.333108
156	1.248461
155	1.189928
98	1.189807

Bottom 10 Pickup/Dropoff Ratios:

	zone	pickup_count	dropoff_count	\
247	Windsor Terrace	27.0	814.0	
0	Newark Airport	198.0	5593.0	
109	Grymes Hill/Clifton	1.0	25.0	
26	Breezy Point/Fort Tilden/Riis Beach	2.0	43.0	
95	Forest Park/Highland Park	3.0	62.0	
122	Inwood Hill Park	6.0	115.0	
44	City Island	3.0	56.0	
191	Ridgewood	60.0	1060.0	
242	Whitestone	23.0	377.0	
192	Riverdale/North Riverdale/Fieldston	42.0	659.0	

	pickup_dropoff_ratio
247	0.033170
0	0.035401
109	0.040000
26	0.046512
95	0.048387
122	0.052174
44	0.053571
191	0.056604
242	0.061008
192	0.063733

3.2.7 [3 marks] Identify zones with high pickup and dropoff traffic during night hours (11PM to 5AM)

```
[652]: # During night hours (11pm to 5am) find the top 10 pickup and dropoff zones
# Note that the top zones should be of night hours and not the overall top zones

# Define night hours (11 PM to 5 AM)
night_hours = list(range(0, 5)) + [23]
# [0, 1, 2, 3, 4, 23] as 4 hour window will have traffic till 4:59 mins
# Filter the DataFrame
night_df = df[df['pickup_hour'].isin(night_hours)]

# Count pickups
night_pickups = night_df.groupby('PULocationID').size().
    ↪reset_index(name='pickup_count')
night_pickups.rename(columns={'PULocationID': 'LocationID'}, inplace=True)
# Count dropoffs
```

```

night_dropoffs = night_df.groupby('DOLocationID').size().
    ↪reset_index(name='dropoff_count')
night_dropoffs.rename(columns={'DOLocationID': 'LocationID'}, inplace=True)

# Merge pickup and dropoff counts
night_traffic = pd.merge(night_pickups, night_dropoffs, on='LocationID',
    ↪how='outer').fillna(0)

# Calculate total traffic
night_traffic['total_night_traffic'] = night_traffic['pickup_count'] +
    ↪night_traffic['dropoff_count']

# Add zone column
night_traffic = night_traffic.merge(zones[['LocationID', 'zone']],
    ↪on='LocationID', how='left')

# Sort by total traffic
top_night_zones = night_traffic.sort_values(by='total_night_traffic',
    ↪ascending=False).head(10)
print(top_night_zones)

top_night_pickup_zones = night_traffic.sort_values(by='pickup_count',
    ↪ascending=False).head(10)
print(top_night_pickup_zones)

top_night_dropoff_zones = night_traffic.sort_values(by='dropoff_count',
    ↪ascending=False).head(10)
print(top_night_dropoff_zones)

```

	LocationID	pickup_count	dropoff_count	total_night_traffic \
78	79	16558.0	8800	25358.0
242	249	13337.0	5190	18527.0
46	48	10464.0	7122	17586.0
127	132	13912.0	1147	15059.0
143	148	10341.0	4629	14970.0
223	230	8308.0	4583	12891.0
67	68	6247.0	5882	12129.0
109	114	9365.0	2672	12037.0
103	107	5778.0	5979	11757.0
159	164	6414.0	4259	10673.0

	zone
78	East Village
242	West Village
46	Clinton East
127	JFK Airport
143	Lower East Side

223	Times Sq/Theatre District				
67	East Chelsea				
109	Greenwich Village South				
103	Gramercy				
159	Midtown South				
	LocationID	pickup_count	dropoff_count	total_night_traffic	\
78	79	16558.0	8800	25358.0	
127	132	13912.0	1147	15059.0	
242	249	13337.0	5190	18527.0	
46	48	10464.0	7122	17586.0	
143	148	10341.0	4629	14970.0	
109	114	9365.0	2672	12037.0	
223	230	8308.0	4583	12891.0	
181	186	6572.0	3698	10270.0	
159	164	6414.0	4259	10673.0	
67	68	6247.0	5882	12129.0	

	zone				
78	East Village				
127	JFK Airport				
242	West Village				
46	Clinton East				
143	Lower East Side				
109	Greenwich Village South				
223	Times Sq/Theatre District				
181	Penn Station/Madison Sq West				
159	Midtown South				
67	East Chelsea				
	LocationID	pickup_count	dropoff_count	total_night_traffic	\
78	79	16558.0	8800	25358.0	
46	48	10464.0	7122	17586.0	
165	170	4173.0	6444	10617.0	
103	107	5778.0	5979	11757.0	
67	68	6247.0	5882	12129.0	
136	141	3408.0	5557	8965.0	
256	263	3442.0	5297	8739.0	
242	249	13337.0	5190	18527.0	
232	239	2745.0	4639	7384.0	
143	148	10341.0	4629	14970.0	

	zone
78	East Village
46	Clinton East
165	Murray Hill
103	Gramercy
67	East Chelsea
136	Lenox Hill West
256	Yorkville West

242 West Village
 232 Upper West Side South
 143 Lower East Side

Now, let us find the revenue share for the night time hours and the day time hours. After this, we will move to deciding a pricing strategy.

3.2.8 [2 marks] Find the revenue share for nighttime and daytime hours.

```
[654]: # Filter for night hours (11 PM to 5 AM)

def get_time_period(hour):
    return 'Night' if hour in [23, 0, 1, 2, 3, 4] else 'Day'

df['time_period'] = df['pickup_hour'].apply(get_time_period)
revenue_by_period = df.groupby('time_period')['total_amount'].sum().
    ↪reset_index()
total_revenue = revenue_by_period['total_amount'].sum()

# Add revenue share column
revenue_by_period['revenue_share'] = revenue_by_period['total_amount'] /
    ↪total_revenue
print(revenue_by_period)
```

	time_period	total_amount	revenue_share
0	Day	50032020.54	0.885585
1	Night	6464009.88	0.114415

Pricing Strategy 3.2.9 [2 marks] For the different passenger counts, find the average fare per mile per passenger.

For instance, suppose the average fare per mile for trips with 3 passengers is 3 USD/mile, then the fare per mile per passenger will be 1 USD/mile.

```
[656]: # Analyse the fare per mile per passenger for different passenger counts

df_nonzerotrip = df[(df['trip_distance'] > 0) & (df['passenger_count'] > 0)]
df_nonzerotrip['fare_per_mile_per_passenger'] = df_nonzerotrip['total_amount'] /
    ↪(df_nonzerotrip['trip_distance'] * df_nonzerotrip['passenger_count'])

fare_stats = df_nonzerotrip.
    ↪groupby('passenger_count')['fare_per_mile_per_passenger'].agg(['mean',
    ↪'median', 'count']).reset_index()
print(fare_stats)
```

C:\Users\Sherry\AppData\Local\Temp\ipykernel_11388\1008790237.py:4:
 SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_nonzerotrip['fare_per_mile_per_passenger'] = df_nonzerotrip['total_amount']  
/ (df_nonzerotrip['trip_distance'] * df_nonzerotrip['passenger_count'])
```

	passenger_count	mean	median	count
0	1.0	16.642746	11.394958	1465055
1	2.0	9.308108	5.486631	288161
2	3.0	5.748642	3.705356	71576
3	4.0	5.848518	2.708333	39462
4	5.0	2.761556	2.245989	24882
5	6.0	2.214246	1.891815	16616

3.2.10 [3 marks] Find the average fare per mile by hours of the day and by days of the week

```
[668]: # Compare the average fare per mile for different days and for different times  
↳ of the day
```

```
# Calculate fare per mile
```

```
df_nonzerotrip['fare_per_mile'] = df_nonzerotrip['total_amount'] /  
↳ df_nonzerotrip['trip_distance']
```

```
fare_by_day = df_nonzerotrip.groupby('pickup_day')['fare_per_mile'].mean().  
↳ reindex([  
    'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday'  
])
```

```
fare_by_hour = df_nonzerotrip.groupby('pickup_hour')['fare_per_mile'].mean()
```

```
print("Fare by Day", fare_by_day)  
print("Fare by hour", fare_by_hour)
```

C:\Users\Sherry\AppData\Local\Temp\ipykernel_11388\3179039869.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_nonzerotrip['fare_per_mile'] = df_nonzerotrip['total_amount'] /  
df_nonzerotrip['trip_distance']
```

Fare by Day pickup_day

Monday	16.690505
Tuesday	17.435274
Wednesday	17.045675
Thursday	17.121152
Friday	16.881214
Saturday	16.122162

```

Sunday          18.050026
Name: fare_per_mile, dtype: float64
Fare by hour pickup_hour
0      15.745927
1      17.019776
2      15.368973
3      16.752811
4      18.852770
5      20.698345
6      15.582748
7      14.966662
8      15.315742
9      15.502052
10     15.781852
11     16.267838
12     17.185356
13     17.479966
14     16.763815
15     18.039635
16     21.347213
17     19.408165
18     18.569847
19     18.463796
20     15.037497
21     14.701038
22     15.384210
23     16.047008
Name: fare_per_mile, dtype: float64

```

3.2.11 [3 marks] Analyse the average fare per mile for the different vendors for different hours of the day

```

[674]: # Compare fare per mile for different vendors
df_nonzerotrip

vendor_fares = df_nonzerotrip.groupby('VendorID')['fare_per_mile'].agg(['mean',
↪ 'median', 'count']).reset_index()
print(vendor_fares)

```

	VendorID	mean	median	count
0	1	13.308087	11.407407	476121
1	2	18.279316	11.274725	1429624
2	6	12.064572	9.900990	7

3.2.12 [5 marks] Compare the fare rates of the different vendors in a tiered fashion. Analyse the average fare per mile for distances upto 2 miles. Analyse the fare per mile for distances from 2 to 5 miles. And then for distances more than 5 miles.

```
[678]: # Defining distance tiers
def distance_tier(dist):
    if dist <= 2:
        return '0-2 miles'
    elif dist <= 5:
        return '2-5 miles'
    else:
        return '5+ miles'

df_nonzerotrip['distance_tier'] = df_nonzerotrip['trip_distance'].
    ↪apply(distance_tier)
tiered_fares = df_nonzerotrip.groupby(['VendorID',
    ↪'distance_tier'])['fare_per_mile'].mean().reset_index()
print (tiered_fares)
plt.figure(figsize=(10, 6))
sns.barplot(data=tiered_fares, x='distance_tier', y='fare_per_mile',
    ↪hue='VendorID', palette='Set2')

plt.title("Average Fare per Mile by Vendor and Distance Tier")
plt.xlabel("Distance Tier")
plt.ylabel("Fare per Mile ($)")
plt.grid(True)
plt.tight_layout()
plt.show()
```

C:\Users\Sherry\AppData\Local\Temp\ipykernel_11388\3727777382.py:10:

SettingWithCopyWarning:

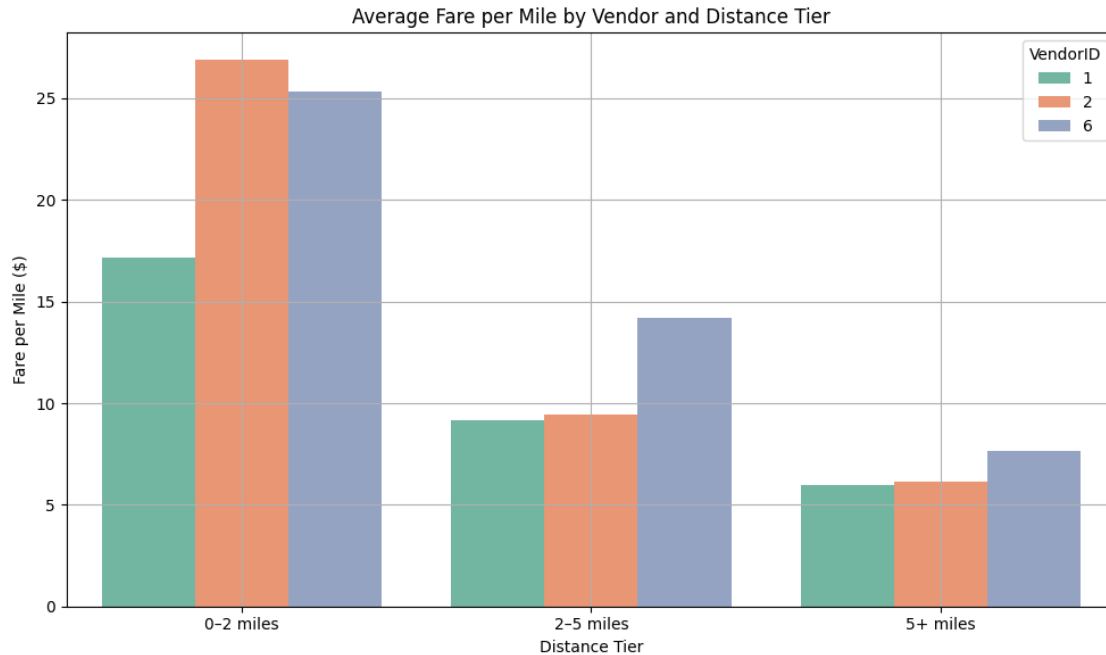
A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_nonzerotrip['distance_tier'] =
df_nonzerotrip['trip_distance'].apply(distance_tier)
```

	VendorID	distance_tier	fare_per_mile
0	1	0-2 miles	17.187697
1	1	2-5 miles	9.184617
2	1	5+ miles	5.990546
3	2	0-2 miles	26.878538
4	2	2-5 miles	9.460546
5	2	5+ miles	6.155048
6	6	0-2 miles	25.333333
7	6	2-5 miles	14.214047
8	6	5+ miles	7.672644



Customer Experience and Other Factors 3.2.13 [5 marks] Analyse average tip percentages based on trip distances, passenger counts and time of pickup. What factors lead to low tip percentages?

```
[686]: # Analyze tip percentages based on distances, passenger counts and pickup times
df['tip_percent'] = (df['tip_amount'] / df['total_amount']) * 100
df['distance_tier'] = df['trip_distance'].apply(distance_tier)
tip_by_distance = df.groupby('distance_tier')['tip_percent'].mean().
    ↪reset_index()
print(tip_by_distance)
tip_by_passenger = df[df['passenger_count'] > 0].
    ↪groupby('passenger_count')['tip_percent'].mean().reset_index()
print(tip_by_passenger)
tip_by_hour = df.groupby('pickup_hour')['tip_percent'].mean().reset_index()
print(tip_by_hour)
```

	distance_tier	tip_percent
0	0-2 miles	12.053035
1	2-5 miles	12.234527
2	5+ miles	11.280166

	passenger_count	tip_percent
0	1.0	12.082660
1	2.0	11.771686
2	3.0	11.358086
3	4.0	10.466442
4	5.0	12.116323

5	6.0	12.209256
	pickup_hour	tip_percent
0	0	11.861455
1	1	11.832382
2	2	11.639764
3	3	11.179685
4	4	10.346102
5	5	10.136287
6	6	11.023918
7	7	11.850445
8	8	12.181061
9	9	12.033606
10	10	11.762246
11	11	11.756538
12	12	11.761377
13	13	11.705225
14	14	11.766097
15	15	11.749014
16	16	11.774628
17	17	12.041503
18	18	12.224745
19	19	12.194702
20	20	12.404647
21	21	12.545225
22	22	12.461128
23	23	12.092979

```
[688]: # Distance
plt.figure(figsize=(10, 5))
sns.barplot(data=tip_by_distance, x='distance_tier', y='tip_percent',
            palette='Blues')
plt.title("Average Tip % by Trip Distance")
plt.ylabel("Tip Percentage")
plt.xlabel("Distance Bin")
plt.tight_layout()
plt.show()


# Passenger Count
plt.figure(figsize=(10, 5))
sns.barplot(data=tip_by_passenger, x='passenger_count', y='tip_percent',
            palette='Greens')
plt.title("Average Tip % by Passenger Count")
plt.ylabel("Tip Percentage")
plt.xlabel("Passenger Count")
plt.tight_layout()
plt.show()
```

```
# Pickup Hour
plt.figure(figsize=(10, 5))
sns.lineplot(data=tip_by_hour, x='pickup_hour', y='tip_percent', marker='o')
plt.title("Average Tip % by Pickup Hour")
plt.xlabel("Hour of Day")
plt.ylabel("Tip Percentage")
plt.grid(True)
plt.tight_layout()
plt.show()
```

C:\Users\Sherry\AppData\Local\Temp\ipykernel_11388\895667400.py:3:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.



```
sns.barplot(data=tip_by_distance, x='distance_tier', y='tip_percent',
palette='Blues')
```



C:\Users\Sherry\AppData\Local\Temp\ipykernel_11388\895667400.py:12:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(data=tip_by_passenger, x='passenger_count', y='tip_percent',
palette='Greens')
```



Additional analysis [optional]: Let's try comparing cases of low tips with cases of high tips to find out if we find a clear aspect that drives up the tipping behaviours

```
[712]: # Count drop-offs by DOLocationID
dropoff_counts = df['DOLocationID'].value_counts()

# Get the location with the maximum drop-offs
max_dropoff_location = dropoff_counts.idxmax()
max_dropoff_count = dropoff_counts.max()
print(dropoff_counts)
```

```

print(f"Location ID with most drop-offs: {max_dropoff_location}\n
↳({max_dropoff_count:,} trips)")
print("zone is ", zones[zones['LocationID'] == max_dropoff_location]['zone'].
↳values[0])

```

DOLocationID

```

236    86548
237    82371
161    76350
230    59943
170    57654

```

...

```

84      9
187     7
204     6
2       4
99      3

```

Name: count, Length: 260, dtype: int64

Location ID with most drop-offs: 236 (86,548 trips)

zone is Upper East Side North

```

[716]: # Compare trips with tip percentage < 10% to trips with tip percentage > 25%
low_tip = df[df['tip_percent'] < 10]
high_tip = df[df['tip_percent'] > 25]

comparison = pd.DataFrame({
    'Low Tip (<10%)': {
        'Avg Fare': low_tip['total_amount'].median(),
        'Avg Distance': low_tip['trip_distance'].median(),
        'Avg Duration (min)': low_tip['trip_duration'].median(),
        'Avg Tip %': low_tip['tip_percent'].mean(),
        'Median Passengers Count': low_tip['passenger_count'].median(),
        'Maximum Drop off': zones[zones['LocationID'] == \
↳low_tip['DOLocationID'].value_counts().idxmax()]['zone'].values[0],
        'Maximum Pick up': zones[zones['LocationID'] == low_tip['PULocationID'].
↳value_counts().idxmax()]['zone'].values[0],
        'Trip Count': len(low_tip)
    },
    'High Tip (>25%)': {
        'Avg Fare': high_tip['total_amount'].median(),
        'Avg Distance': high_tip['trip_distance'].median(),
        'Avg Duration (min)': high_tip['trip_duration'].median(),
        'Avg Tip %': high_tip['tip_percent'].mean(),
        'Median Passengers Count': high_tip['passenger_count'].median(),
        'Maximum Drop off': zones[zones['LocationID'] == \
↳high_tip['DOLocationID'].value_counts().idxmax()]['zone'].values[0],

```

```

        'Maximum Pick up': zones[zones['LocationID'] == high_tip['PULocationID']].
        ↪value_counts().idxmax()['zone'].values[0],
        'Trip Count': len(high_tip)
    }
})
print(comparison)

```

	Low Tip (<10%)	High Tip (>25%)
Avg Fare	19.6	21.0
Avg Distance	1.8	1.17
Avg Duration (min)	13.1	8.3
Avg Tip %	2.452221	34.486054
Median Passengers Count	1.0	1.0
Maximum Drop off	Upper East Side South	Upper East Side North
Maximum Pick up	JFK Airport	Times Sq/Theatre District
Trip Count	655277	9581

3.2.14 [3 marks] Analyse the variation of passenger count across hours and days of the week.

```

[718]: # See how passenger count varies across hours and days

# Average passenger count by hour and day
passenger_trends = df.groupby(['pickup_day', 'pickup_hour'])['passenger_count'].
    ↪mean().reset_index()

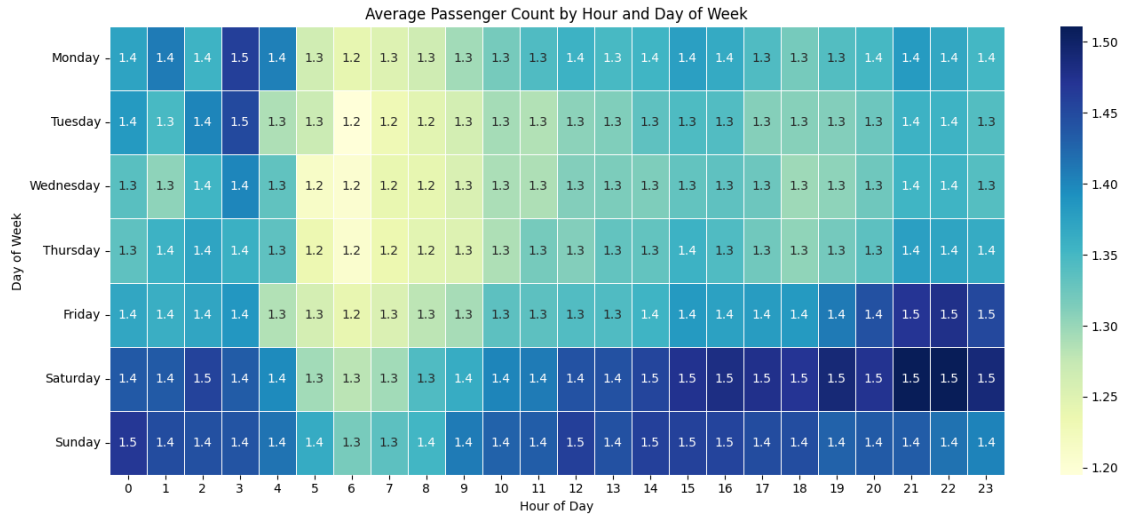
# Pivot for heatmap
pivot_table = passenger_trends.pivot(index='pickup_day', columns='pickup_hour',
    ↪values='passenger_count')

ordered_days = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
    ↪'Saturday', 'Sunday']
pivot_table = pivot_table.reindex(ordered_days)

# Plot
plt.figure(figsize=(14, 6))
sns.heatmap(pivot_table, cmap='YlGnBu', annot=True, fmt=".1f", linewidths=0.5)

plt.title("Average Passenger Count by Hour and Day of Week")
plt.xlabel("Hour of Day")
plt.ylabel("Day of Week")
plt.tight_layout()
plt.show()

```



3.2.15 [2 marks] Analyse the variation of passenger counts across zones

[720]: # How does passenger count vary across zones

```
zone_passenger_stats = df.groupby('zone')['passenger_count'].agg(['mean',
    ↪ 'median', 'sum', 'count']).reset_index()
zone_passenger_stats = zone_passenger_stats.sort_values(by='mean',
    ↪ ascending=False)

top_zones = zone_passenger_stats.head(10)

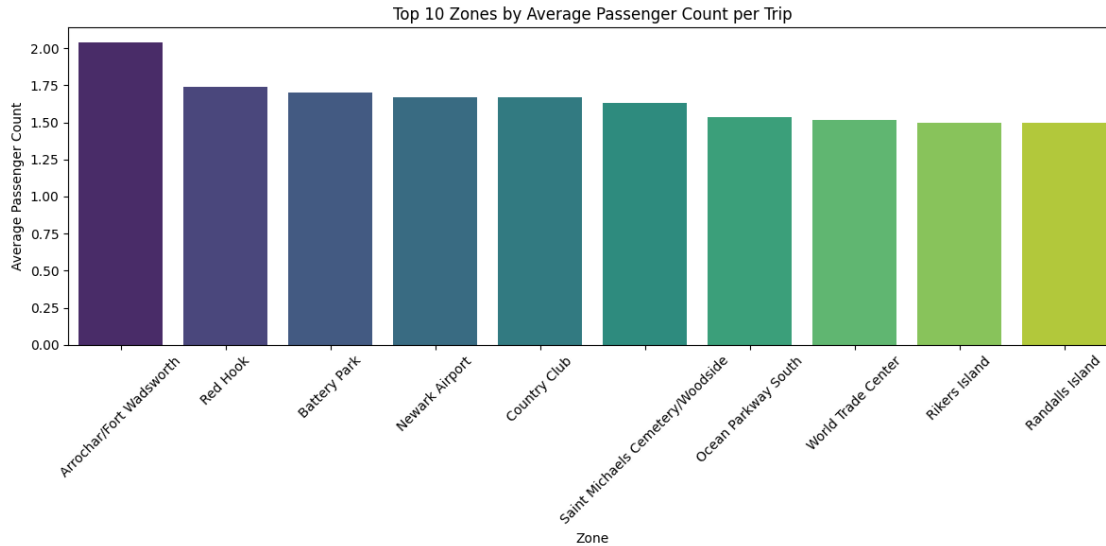
plt.figure(figsize=(12, 6))
sns.barplot(data=top_zones, x='zone', y='mean', palette='viridis')
plt.title("Top 10 Zones by Average Passenger Count per Trip")
plt.ylabel("Average Passenger Count")
plt.xlabel("Zone")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

C:\Users\Sherry\AppData\Local\Temp\ipykernel_11388\4228150114.py:9:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(data=top_zones, x='zone', y='mean', palette='viridis')
```



```
[724]: # For a more detailed analysis, we can use the zones_with_trips GeoDataFrame
# Create a new column for the average passenger count in each zone.
avg_passenger = df.groupby('zone')['passenger_count'].mean().
    ↪reset_index(name='avg_passenger_count')
zones_with_trips = zones.merge(avg_passenger, on='zone', how='left')
```

```
[724]:
```

	OBJECTID	Shape_Leng	Shape_Area	zone	LocationID	\
0	1	0.116357	0.000782	Newark Airport	1	
1	2	0.433470	0.004866	Jamaica Bay	2	
2	3	0.084341	0.000314	Allerton/Pelham Gardens	3	
3	4	0.043567	0.000112	Alphabet City	4	
4	5	0.092146	0.000498	Arden Heights	5	
..	
258	259	0.126750	0.000395	Woodlawn/Wakefield	259	
259	260	0.133514	0.000422	Woodside	260	
260	261	0.027120	0.000034	World Trade Center	261	
261	262	0.049064	0.000122	Yorkville East	262	
262	263	0.037017	0.000066	Yorkville West	263	

	borough	geometry	\
0	EWB	POLYGON ((933100.918 192536.086, 933091.011 19...	
1	Queens	MULTIPOLYGON (((1033269.244 172126.008, 103343...	
2	Bronx	POLYGON ((1026308.77 256767.698, 1026495.593 2...	
3	Manhattan	POLYGON ((992073.467 203714.076, 992068.667 20...	
4	Staten Island	POLYGON ((935843.31 144283.336, 936046.565 144...	
..	
258	Bronx	POLYGON ((1025414.782 270986.139, 1025138.624 ...	
259	Queens	POLYGON ((1011466.966 216463.005, 1011545.889 ...	

```

260     Manhattan POLYGON ((980555.204 196138.486, 980570.792 19...
261     Manhattan MULTIPOLYGON (((999804.795 224498.527, 999824...
262     Manhattan POLYGON ((997493.323 220912.386, 997355.264 22...

```

	pickup_count	dropoff_count	total_trips	avg_passenger_count
0	198.0	5593.0	5791.0	1.666667
1	2.0	4.0	6.0	1.000000
2	39.0	164.0	203.0	1.025641
3	2273.0	7370.0	9643.0	1.354158
4	13.0	34.0	47.0	1.000000
..
258	48.0	212.0	260.0	1.104167
259	328.0	1398.0	1726.0	1.290520
260	10289.0	9280.0	19569.0	1.514677
261	25874.0	30149.0	56023.0	1.299126
262	37490.0	40052.0	77542.0	1.312758

[263 rows x 11 columns]

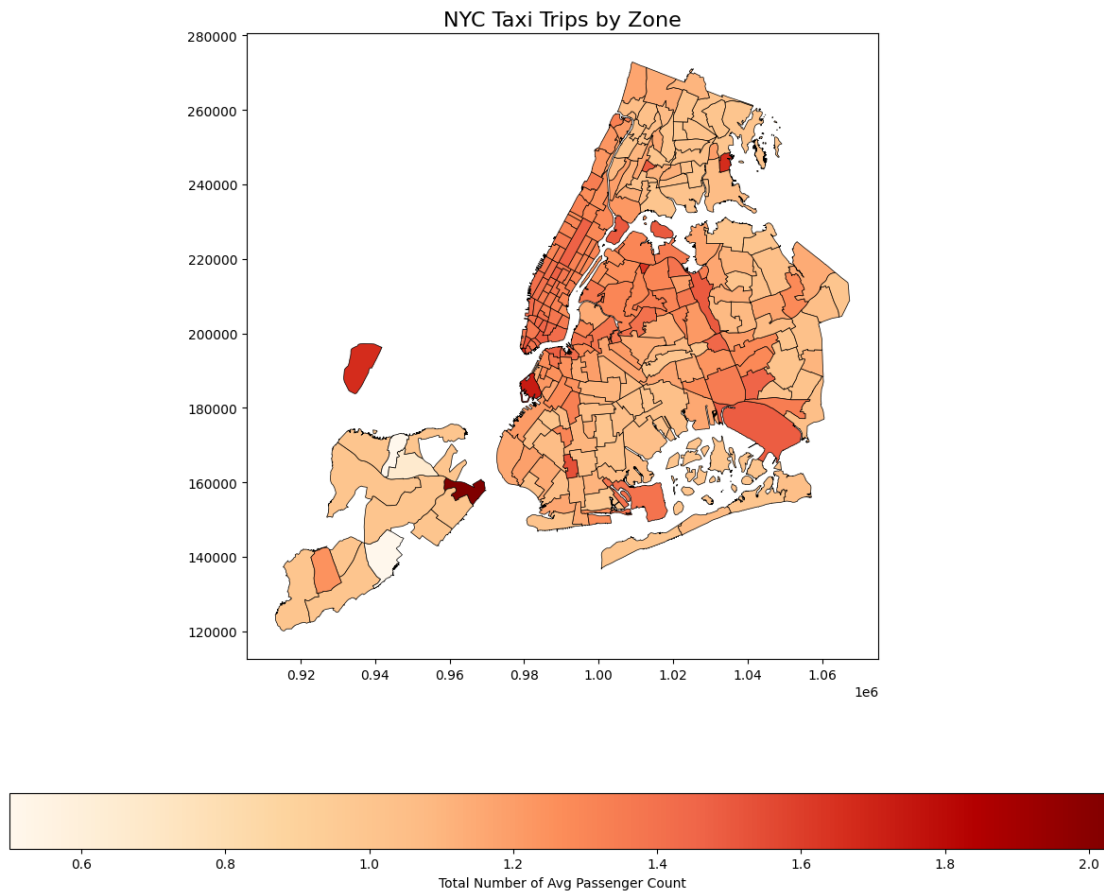
```

[728]: # Define figure and axis
fig, ax = plt.subplots(1, 1, figsize=(12, 10))

# Plot the map and display it
zones_with_trips.plot(column='avg_passenger_count',
                      cmap='OrRd',
                      legend=True,
                      edgecolor='black',
                      linewidth=0.5,
                      ax=ax,
                      legend_kwds={
                          'label': "Total Number of Avg Passenger Count",
                          'orientation': "horizontal" # or "vertical"
                      })

ax.set_title("NYC Taxi Trips by Zone", fontsize=16)
plt.tight_layout()
plt.show()

```

Find out how often surcharges/extra charges are applied to understand their prevalence

```
[730]: surcharge_cols = ['extra', 'mta_tax', 'improvement_surcharge',
    ↪ 'congestion_surcharge']
```

```
# Count how often each surcharge is applied (non-zero)
surcharge_stats = {
    col: {
        'applied_count': (df[col] > 0).sum(),
        'total_trips': len(df),
        'prevalence_percent': 100 * (df[col] > 0).mean()
    }
    for col in surcharge_cols
}
```

```
surcharge_df = pd.DataFrame(surcharge_stats).T
print(surcharge_df)
```

```
applied_count  total_trips  prevalence_percent
```

extra	1188170.0	1947260.0	61.017532
mta_tax	1931825.0	1947260.0	99.207348
improvement_surcharge	1947049.0	1947260.0	99.989164
congestion_surcharge	1807575.0	1947260.0	92.826587

3.2.16 [5 marks] Analyse the pickup/dropoff zones or times when extra charges are applied more frequently

```
[732]: df.rename(columns={'zone': 'pickup_zone'}, inplace=True)
```

```
[734]: # Merge dropoff zone names into df
df = df.merge(zones[['LocationID', 'zone']], left_on='DOLocationID',
             right_on='LocationID', how='left')

# Rename the merged column to 'dropoff_zone'
df.rename(columns={'zone': 'dropoff_zone'}, inplace=True)
```

```
-----
KeyError                                Traceback (most recent call last)
Cell In[734], line 8
      5 df.rename(columns={'zone': 'dropoff_zone'}, inplace=True)
      7 # Optionally drop the extra LocationID column if not needed
----> 8 df.drop(columns='LocationID', inplace=True)

File ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\frame.py:5581
in DataFrame.drop(self, labels, axis, index, columns, level, inplace, errors)
    5433 def drop(
    5434     self,
    5435     labels: IndexLabel | None = None,
    (...)
    5442     errors: IgnoreRaise = "raise",
    5443 ) -> DataFrame | None:
    5444     """
    5445     Drop specified labels from rows or columns.
    5446
    (...)
    5579         weight 1.0      0.8
    5580     """
-> 5581     return super().drop(
    5582         labels=labels,
    5583         axis=axis,
    5584         index=index,
    5585         columns=columns,
    5586         level=level,
    5587         inplace=inplace,
    5588         errors=errors,
    5589     )
```

```

File ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\generic.py:
  4788, in NDFrame.drop(self, labels, axis, index, columns, level, inplace,
  4789 errors)
    4786 for axis, labels in axes.items():
    4787     if labels is not None:
-> 4788         obj = obj._drop_axis(labels, axis, level=level, errors=errors)
    4790 if inplace:
    4791     self._update_inplace(obj)

File ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\generic.py:
  4830, in NDFrame._drop_axis(self, labels, axis, level, errors, only_slice)
    4828     new_axis = axis.drop(labels, level=level, errors=errors)
    4829     else:
-> 4830         new_axis = axis.drop(labels, errors=errors)
    4831     indexer = axis.get_indexer(new_axis)
    4833 # Case for non-unique axis
    4834 else:

File ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\indexes\base.
  py:7070, in Index.drop(self, labels, errors)
    7068 if mask.any():
    7069     if errors != "ignore":
-> 7070         raise KeyError(f"{labels[mask].tolist()} not found in axis")
    7071     indexer = indexer[~mask]
    7072 return self.delete(indexer)

KeyError: "['LocationID'] not found in axis"

```

```
[738]: df.drop(columns='LocationID_y', inplace=True)
```



```
[740]: # How often is each surcharge applied?
df['has_surcharge'] = df[surcharge_cols].sum(axis=1) > 0
# Pickup zones
pickup_surcharge = df.groupby('pickup_zone')['has_surcharge'].mean().
    reset_index(name='pickup_surcharge_rate')

# Dropoff zones (if you have dropoff zone names)
dropoff_surcharge = df.groupby('dropoff_zone')['has_surcharge'].mean().
    reset_index(name='dropoff_surcharge_rate')
```

```
[742]: hourly_surcharge = df.groupby('pickup_hour')['has_surcharge'].mean().
    reset_index(name='surcharge_rate')
```

```
[748]: # By hour
plt.figure(figsize=(10, 5))
sns.lineplot(data=hourly_surcharge, x='pickup_hour', y='surcharge_rate',
    marker='o')
```

```
plt.title("Surcharge Rate by Hour of Day")
plt.xlabel("Hour")
plt.ylabel("Surcharge Rate")
plt.grid(True)
plt.tight_layout()
plt.show()
```



1.8 4 Conclusion

[15 marks]

1.8.1 4.1 Final Insights and Recommendations

[15 marks]

Conclude your analyses here. Include all the outcomes you found based on the analysis.

Based on the insights, frame a concluding story explaining suitable parameters such as location, time of the day, day of the week etc. to be kept in mind while devising a strategy to meet customer demand and optimise supply.

4.1.1 [5 marks] Recommendations to optimize routing and dispatching based on demand patterns and operational inefficiencies

- []:
- Plan **for** taxis to be available between 2 PM **and** 7 PM on weekdays, when demand **is** typically at its highest.
 - For nighttime operations, schedule taxis between 12 AM **and** 5 AM on weekends, **as** demand tends to be significantly higher during these hours.
 - When planning **for** holidays, prioritize either Monday **or** Sunday, **as** pickup **volume is** generally lower on these days.
 - May **and** October usually see a surge **in** pickup demand, so plan fleet **availability** accordingly.
 - Focus on shorter trips lasting less than 2 to 3 hours to maximize turnaround **and** efficiency.
 - Since most customers prefer paying by credit card, ensure that taxis are **equipped with** card payment facilities.

4.1.2 [5 marks]

Suggestions on strategically positioning cabs across different zones to make best use of insights uncovered by analysing trip trends across time, days and months.

- []:
- Position your taxis **in** high-demand zones such **as** Upper East Side South, Upper **East Side North, and** Midtown Center, where both pickups **and** drop-offs are **consistently high**.
 - Avoid scheduling pickups **and** drop-offs on slower routes based on time-of-day **variations**. This insight can be communicated to drivers to help them choose **the fastest available routes**.
 - Since demand peaks between 5 PM **and** 7 PM, plan **for** efficient routing during **this window to improve turnaround time**.
 - By analyzing hourly pickup **and** drop-off trends, you can better estimate the **likelihood of securing a fare in** specific areas.
 - Certain zones **like** East Elmhurst, JFK Airport, **and** LaGuardia Airport **show high pickup-to-drop-off ratios**. Taxis should prioritize trips to **and from** **these** areas.
 - Based on nighttime pickup **and** drop-off patterns, taxis can strategically plan **their routes to maximize efficiency during late hours**

4.1.3 [5 marks] Propose data-driven adjustments to the pricing strategy to maximize revenue while maintaining competitive rates with other vendors.

- []:
- Based on hourly pickup and drop-off zone trends, dynamic pricing strategies can be developed.
 - During high-demand periods—especially between 5 PM and 7 PM—fares can be adjusted upward to optimize revenue.
 - To maximize earnings, prioritize daytime operations over nighttime, as demand is typically higher during the day.
 - Fare pricing should reflect peak demand at specific hours throughout the day.
 - Vendor 1 may increase rates for both shorter and longer trips.
 - Vendor 2 can focus fare adjustments on trips longer than 2 miles.
 - Revenue is generally higher during March, April, May, June, October, November, and December compared to other months.
 - If planning vehicle servicing or holidays, February is ideal, as it tends to have the lowest revenue. Further analysis of pickup times and zones may help boost revenue in off-peak months.
 - Trips under 50 miles tend to be more profitable. Avoid assigning taxis to longer routes, as longer duration doesn't always mean higher earnings.
 - Tips are typically higher for shorter trips, whereas between 3 AM and 7 AM, tip amounts are significantly lower.
 - Surcharges are least common from 5 AM to 9 AM, which could impact fare planning during early morning hours.